

Pitstop Operations Manual

The self-serve guide to customer support tickets, calls and admin at Legend of Toys.



Version
1.22.1

Updated
2026-09-04

Application
pitstop.legendoftoys.com

Contents

GETTING STARTED

Welcome to Pitstop	4
Signing In & Getting Around	6
Roles & What You Can See	7

WORK

The Queue /queue	10
Working a Ticket /queue/detail	12
New Ticket /new	15
The Inbox /inbox	16
Post Comments /comments	22
Calls /calls	24
Call Detail /calls/detail	27

ANALYZE

Support Analytics /analytics	29
Reports /reports	32

ADMIN

Departments /admin/departments	38
MyOperator Accounts /admin/myop	40
Telephony /admin/telephony	42
WhatsApp Templates /admin/wa-templates	44
Meta Diagnostics /admin/meta	45

PART 1

Getting Started

What Pitstop is, how to sign in and get around, and what your role can see.

Welcome to Pitstop

AGENT

LEAD

ADMIN

Pitstop is the customer support system for Legend of Toys. Every customer issue, whether it arrives by phone, WhatsApp or a marketplace message, becomes a **ticket** that an agent works from start to finish. This manual is your self-serve guide to every screen in it.

How a ticket flows

Almost everything in Pitstop revolves around a ticket moving through clear stages. The shape of that journey is set by one choice: the **disposition**, which is what you have decided to do for the customer.

DISPOSITION	WHAT HAPPENS
Query / No action	A question or a non-issue: opens and closes quickly (intake to closed).
Awaiting info	Waiting on the customer for evidence before deciding.
Replacement / Refund / Repair / Pending	The full path: verify, schedule a pickup, receive at the warehouse, inspect, then the branch for the chosen outcome, then close.

You set the disposition during triage; from there the ticket advances stage by stage, each step asking for the details it needs (an AWB, an inspection note, a refund reference).

① WHERE THE CONVERSATION HAPPENS

Every customer conversation lives in Pitstop, in the **Inbox**: WhatsApp, Instagram, Messenger, email and the website chat, all two-way. A ticket shows the conversation it came from as a read-only panel with an **Open in inbox** button, so you read it beside the case and reply in the one place that sends. Phone calls flow in from Exotel and appear in the Calls log, where a missed call can be turned into a ticket.

How to use this manual

Each chapter covers one screen and follows the same shape: what it is for, how to read it, and what to do. Jump to the screen you are on using the Contents page or your PDF reader's bookmarks. Chapters are tagged with the roles that use them:

AGENT

LEAD

ADMIN

Callout boxes

✓ **TIP**

A faster or smarter way to do something.

⚠ **WATCH OUT**

A common mistake or something easy to misread.

⚡ **IMPORTANT**

Something that affects a customer outcome or money: get it right.

Signing In & Getting Around



Pitstop opens with one click using your work Google account, and everything is reachable from a short sidebar. This chapter covers signing in and finding your way around.

Signing in

Open `pitstop.legendoftoys.com` and sign in with your `@legendoftoys.com` Google account. Personal accounts are refused; if sign-in bounces you back, sign out of any personal Google account first or use a fresh browser profile.

The menu

GROUP	WHAT LIVES HERE
WORK	Overview, Queue, Inbox, Calls, New Ticket.
ANALYZE	Reports (for leads and admins).
ADMIN	Departments, MyOperator Accounts, WhatsApp Templates (admins only).

You only see the groups your role allows. Most screens have a search box you can focus by pressing `/`, and the queue and calls lists support arrow keys with Enter to open a row. **On a phone** the side rail is replaced by the tab bar at the bottom. Tap the **magnifying-glass icon in the bar across the top** to open the **command palette** — type a screen name and press Enter. That is a different thing from the `/` box above, which filters only the list you are looking at.

The sidebar **collapses** to a slim strip of icons to give you more room, and starts collapsed by default. Click the chevron at the bottom of the sidebar, or click the **PITSTOP** header, to expand or collapse it; it remembers your choice. When collapsed, hover an icon to see its label.

① THE DEPARTMENT SWITCHER

Admins get a **department switcher** in the top bar that scopes every list to one team or “All Departments”. Everyone else is locked to their own department automatically, so you only see your team’s tickets and calls. Customer phone numbers are **masked** by default across the app, with a reveal toggle only on the ticket’s customer panel.

Roles & What You Can See



Pitstop gives each person the access that matches their job. This chapter explains the three roles used throughout this manual and what each unlocks.

The three roles

AGENT	CS Agent. Works the queue and calls: opens tickets, verifies issues, advances them through their stages, logs notes and evidence, and claims tickets to themselves.
LEAD	CS Lead. Everything an agent does, plus reassigning tickets to other agents and reading the Reports.
ADMIN	Admin. Full access: changing a ticket’s disposition after triage, the department switcher, and the admin screens (Departments, MyOperator accounts, WhatsApp templates).

① HOW ACCESS ACTUALLY WORKS

Your access comes from the permissions on your account (for example `cs_ticket_manage` to work tickets, `cs_ticket_reassign` to assign to others, `cs_reports_view` for reports, `cs_ticket_admin` for the admin screens). The three tags above are plain-language groupings of those permissions. If you cannot do something you should be able to, an admin needs to adjust your role.

Who sees what

AREA	AGENT	LEAD	ADMIN
Queue, Calls, New Ticket, working a ticket	•	•	•
Reassign tickets to other agents		•	•
Reports		•	•
Change disposition past triage			•

AREA	AGENT	LEAD	ADMIN
Departments, MyOp, WA Templates, dept switcher			•

PART 2

Work

The queue, the ticket workspace, creating tickets, the conversations inbox, and the call log.

The Queue

/queue

AGENT

LEAD

ADMIN

The Queue is the CS home screen: the worklist of tickets. Agents live here to see what is assigned to them, what is overdue, and to open or create tickets. It is where a shift starts.

WHAT IT'S FOR	READ OR ACT?	WHO	UPDATES
Your ticket worklist.	Read; open or create tickets.	AGENT LEAD ADMIN	Counts every 60s.

Reading it

Five cards head the screen: **My Open, Overdue** (red), **Awaiting Evidence 3+ days** (nudge candidates), **Closed Today**, and **Avg Close** month-to-date. Below them, tabs with live counts: **My Queue** (the default), All Open, Awaiting Evidence, In Logistics, Inspection, Resolution, **Closure Requests**, Closed.

Closure Requests lists every open ticket an agent has asked to close that is still waiting for an admin's decision. Agents don't close tickets directly — the Close button on a ticket submits a *request* with an outcome and a compulsory note, and the ticket stays where it is until an admin approves or rejects it from the banner on the ticket page. Anywhere in the queue, a pending request also shows as a small amber `closure req` marker under the ticket's stage, naming who asked and the outcome they chose.

Each row shows the ticket number, customer (with masked phone), product, the **disposition** badge (an `AUTO` tag if it was auto-created), the current stage, platform, assigned agent, and an **Age / SLA** figure: green up to 3 days, orange beyond, and **red and bold when overdue** with a "+Nd over" note. The exact **creation date and time** sits in small grey type under the age (and on hover), shown in Indian time so you can tell at a glance when a ticket actually came in.

Finding and opening tickets

Search by name, phone, order, UPC or ticket number (press `/` to jump in). Narrow with the disposition chips and the Category and Platform dropdowns; the filters live in the URL so a view can be bookmarked. Click a row, or use the arrow keys and Enter, to open it. **New Ticket** (red) creates one; **Export CSV** downloads the rows you can currently see.

Every tab shows 50 tickets at a time. Use the **Prev** and **Next** buttons at the bottom to page through the rest; the footer shows which rows and page you are on. Changing the tab, any filter or the search starts you back at the first page.

⚠ EXPORT IS THE VISIBLE LIST, WITH MASKED PHONES

Export CSV downloads only the tickets currently loaded for the tab and filters you are on, not the whole queue, and phone numbers are masked in the file. The `AUTO` badge means the ticket was created automatically (for example from a call), not opened by an agent.

Working a Ticket

/queue/detail

AGENT

LEAD

ADMIN

The ticket detail screen is the full case file for one ticket. This is where you verify the issue, advance it through its lifecycle, read the customer conversation, log notes and evidence, and assign it. It is the heart of Pitstop.

WHAT IT'S FOR	READ OR ACT?	WHO	LAYOUT
Working a single ticket end to end.	The main action screen.	AGENT LEAD ADMIN	Identity · Work area · Activity.

The layout

A header carries the customer and product (with its colour where known), the disposition, an overdue badge if relevant, the creation date and time in Indian time, and the big red **Advance** button, with a stage **stepper** below. The body has three columns: the **identity rail** (assignment, customer, Shopify, call info, order, past cases), the **work area** (Issue, Return logistics, Resolution), and the **activity feed** (history, notes, attachments).

The **Shopify panel** on the identity rail pulls the customer's full order history: each recent order lists its line items (product, quantity, variant and price), the order total, where it shipped, and the courier tracking number with a **track** link, so you can confirm what was bought and where it is without leaving the ticket.

Advancing a ticket

- 1

Set the **disposition** from the header dropdown during triage. This decides the whole path the ticket will take.
- 2

Click **Advance**. The button names the next stage (for example "Schedule Pickup"). A modal asks for the fields that stage needs: a pickup needs the return AWB and courier; inspection needs an inspection note; a refund needs the amount, then the payment reference. **Verifying the issue is the exception**: it has no fields, but it needs at least one evidence attachment first: the modal tells you so and how to add one.
- 3

Fill the **required fields and confirm**. The ticket moves to the next stage and the stepper updates.

Close now sits beside Advance for admins, and it is the shortcut for a ticket that is finished long before its stage ladder is. A replacement ticket has eight rungs between intake and closed, each with its own required fields — when the case is simply over (the customer stopped needing it, it turned out to be a

duplicate, it was handled on the phone), you should not have to walk it through pickup, warehouse and inspection to end it. Click **Close now**, pick the outcome from the same list the normal close uses, add a note if it helps, and the ticket closes where it stands with its disposition intact.

Close now is an admin action, and deliberately so. It skips every remaining stage gate, so only an admin sees it. It also does not appear when Advance is already the close button (there would be two buttons doing one job), nor while a closure request is waiting — approve or reject that from the banner instead. Agents who need a ticket closed mid-way should ask an admin; the Request Closure route only opens at the last rung.

Three side-exits handle the non-standard endings: **Cancel** (customer dropped), **Reject** (inspection found no fault), and **Escalate** (flag for a manager, staying in the current stage). Assignment is on the left rail: agents **Claim** an unassigned ticket; leads and admins can **Reassign** to anyone.

▲ DISPOSITION LOCKS AFTER TRIAGE

You can change the disposition freely only while the ticket is in **intake** or **awaiting evidence**. Once it enters logistics, only an admin can change it. The Issue, Return logistics and Resolution panels also unlock in sequence: a panel that is greyed out now ("unlocks once the issue is verified") will open at the right stage. Cancel, Reject and Close are three different endings: do not use one for another.

Switching between Replacement and Refund

A ticket that was set to **Replacement** or **Refund** can be flipped to the other one with the ⇌ **Refund** / ⇌ **Replacement** button next to the disposition in the header. Any agent can do this, and you are asked for an optional reason that is saved to the ticket history. This is the one disposition change that does not need an admin after triage.

▲ SWITCH ONLY BEFORE THE RESOLUTION STARTS

The switch works only while the ticket has not yet started its resolution: that is, before a refund has been initiated or a replacement dispatched. Once the refund or replacement is under way (or the ticket is closed), the button disappears: at that point you must **Cancel** and reopen the ticket to change course, so an in-flight refund and a dispatched replacement never run on the same ticket.

① THE CONVERSATION PANEL AND ATTACHMENTS

At the foot of the ticket sits the customer's conversation, on whichever channel raised it: WhatsApp, Instagram, Messenger, email or web chat. It is read only, so you can read the history without leaving the case; to reply, click **Open in inbox**, which opens that exact conversation in the Inbox. A green pill shows how long the 24-hour reply window is open on the channels that have one. Attachments in the activity feed are **links** (to a WhatsApp video, a Drive photo), not uploaded files.

▲ LINKED OR JUST MATCHED

If the panel says **matched by phone number** or **matched by email address**, that conversation is the same customer but is *not* attached to this ticket: it is a lead, not a fact, so read it as background. Tickets raised from a conversation are linked already. To bind one, open it in the Inbox and use **Link ticket**.

New Ticket

/new AGENT LEAD ADMIN

New Ticket logs a customer issue, usually while on a call. It is built so you scan or type the unit's UPC first and the product details fill themselves in.

WHAT IT'S FOR	READ OR ACT?	WHO	REQUIRED
Creating a customer ticket.	Action: creates a ticket.	AGENT LEAD ADMIN	Customer name + description.

Filling it in

- 1 Order lookup.** Scan or type the unit **UPC**; after a few characters it looks the unit up and auto-fills product, model, colour and SKU. If you do not have a UPC, choose from the dropdowns instead. Start with **Category** (L.O.T Cars, L.O.T Build or L.O.T DIY) to narrow the product list, then pick the **Product**, the **Model** and the **Colour**, and the **SKU** fills in by itself. Each list is narrowed by the choice above it. Category is only a filter to help you find the product — it is not saved on the ticket, and you can leave it alone and search the full product list if you prefer. If a product is already chosen, Category shows the one it belongs to. Pick the platform and enter the external order ID if there is one.
- 2 Customer.** Enter the name (required), phone, email and address. When you leave the phone field it checks for past cases and warns you with links if this customer has contacted before.
- 3 Issue.** Pick a category and sub-category (or "Other" with a description), set the disposition (default **Pending**), choose the intake channel, and write the description (required).
- 4 Create Ticket.** The ticket is created and opens on its detail page.

① ONLY TWO FIELDS ARE REQUIRED

Only the customer **name** and the issue **description** are mandatory; the UPC and product can be blank if you do not have them. When you open this from a missed call, the phone and name are pre-filled and the description is seeded for you.

The Inbox



The Inbox is the conversations console: every customer message in one place, across Instagram, Facebook Messenger, WhatsApp and email. Agents read, prioritise, reply, and turn a chat into a ticket without leaving the screen.

WHAT IT'S FOR	READ OR ACT?	WHO	UPDATES
Customer DMs, all channels.	Reply, assign, tag, prioritise, raise tickets.	AGENT LEAD ADMIN	List every 20s, open chat every 15s.

The layout

Tiles across the top show the volume per channel and how many are awaiting a reply; click one to filter to that channel. The **conversation list** sits on the left, the open **conversation** on the right. Three sets of controls narrow the list:

- **Channel** tabs: All, Instagram, Messenger, WhatsApp, Email.
- **Mine / Unassigned / All**: whose conversations to show. Mine is what you own, Unassigned is the pool you can claim from.
- **Active / Closed / All**: Active is open and snoozed chats, Closed is the ones marked done.

Each row shows the customer, the last message preview, time since the last message, and small badges: priority (when not Normal), **Done** if closed, who it is assigned to, a linked ticket number, and any tags.

Filtering and sorting

Under the tabs are a **Sort** control and filters. Sort by **Recent activity** (the default), **Oldest first**, or **Priority** (Urgent at the top). Filter by **priority**, by **tag**, and, for leads and admins, by the **assigned agent**. These stack on top of the channel and status choices, so you can line up, say, all Urgent unanswered Instagram chats.

Priority

Every conversation has a priority you set from the dropdown in the chat header: **Urgent**, **High**, **Normal**, or **Low**. New chats start at Normal. Anything above Normal shows a coloured badge in the list, and you can sort or filter by it, so the urgent ones surface first.

Tags, claiming and transfer

Use the **tag** rail under the header to label a chat (search existing tags or create one). Use **Claim** to take a chat for yourself, **Release** to hand it back to the pool, and **Transfer...** to pass it to a teammate: pick the agent, add an optional note for them, and the chat moves to their queue with your note attached. You can transfer a chat you are handling or an unassigned one; leads and admins can transfer any chat, including one assigned to someone else.

Replying to a customer claims the chat for you, and also assigns the linked ticket (if there is one) to you, so whoever answers owns the case.

① CHATS CAN NOW ARRIVE IN MINE WITHOUT YOU CLAIMING THEM (27 AUG 2026)

Conversations are handed out automatically to whoever is available, and from 27 August that also happens *after the fact*: a chat that arrived while nobody was available used to sit in **Unassigned** forever, and is now picked up and given an owner once someone is back. So you may find chats in **Mine** that you did not claim.

They come in small batches rather than all at once, and only chats with activity in the **last 7 days** are handed out — older ones are left alone deliberately. If one is not right for you, **Release** puts it back in the pool or **Transfer...** passes it to someone specific.

Why it changed: a chat with a name on it gets answered and one in a shared pile does not. Nine hundred Instagram conversations had been sitting unowned, some for over a month.

To move several chats at once, tick the checkbox on each row (or the **Select** box at the top of the list to take them all), then pick a person from the **Assign to...** dropdown and press Apply. You can claim them for yourself, hand them to another agent (leads and admins), or release them back to the pool, all in one go.

A THIRD BANNER: "OUTGOING MESSAGES: SENDING IS DOWN"

Messages the system sends by itself

A few messages in a thread are not from any agent. They are shown as sent by `ReLay (auto)` with no agent name, and there are two of them. The **wrong-number redirect** answers someone who has written to a marketing or transactional number, pointing them at the support number. The **out-of-hours auto-reply** tells a customer who writes when nobody is rostered that the team is offline and will come back the next working day.

A thread gets at most one automated message in any 24 hours, whichever fires first — someone sending five messages overnight is answered once, not five times, and a wrong-number contact arriving at midnight gets the redirect rather than both. The system also stays quiet on any thread an agent has replied to in the last 12 hours, so an automated line never lands on top of a conversation a person is already having. **You still own the thread:** an auto-reply is an acknowledgement, not an answer, and the customer is still waiting for you.

This one is not about a single channel — it means **nothing is leaving the system at all**: replies, automated messages and campaigns alike. It appeared for the first time after 27 August 2026, when outgoing WhatsApp stopped for 52 minutes and nobody noticed until an agent reported it.

What to do: messages coming *in* are unaffected and nothing is being lost, so keep reading and tagging as normal. **Do not keep retyping a reply that shows failed** — it will not go until the banner clears, and it will send once it does. Tell Afshaan; this is not something an agent can fix.

If a whole channel goes quiet

A red banner can appear across the top of the Inbox saying, for example, *"Instagram: nobody is replying"*, with the number of customer messages that arrived in the last 24 hours and how long it has been flagged. It means Pitstop checked and found that customers are writing in on that channel while no agent reply has gone out at all.

This is worth reading carefully, because it catches a problem you cannot see from your own screen. In August 2026 Instagram replies failed for three days and it only came to light when someone mentioned it in Slack. Automatic messages kept going out the whole time, so from the outside the channel looked busy and healthy; it was only the agents' own replies that were being refused. This banner watches for exactly that.

① WHAT TO DO, AND WHY YOU CANNOT DISMISS IT

There is no close button, on purpose: it clears by itself once replies start going out again, so it cannot be waved away and forgotten while customers are still waiting. If you see it, first try sending one reply on that channel. If your reply fails, report it straight away and say which channel; if your reply works, the banner will disappear on its own within a few minutes.

It is checked four times a day rather than continuously, so a channel that has just broken may take a few hours to be flagged. Treat the banner as confirmation, never as your only warning.

Alerts: badge, chime and desktop notifications

Three things tell you a new customer message has arrived, and they are deliberately different so that at least one reaches you wherever you are.

- **The tab badge** is always on. The browser tab title shows the number awaiting a reply, so a glance at the tab tells you without switching to it.
- **The chime** is the bell button above the conversation list. Click it to turn a short two-tone sound on or off when a new message arrives. Your choice is remembered on that computer.
- **Desktop notifications** are the screen button beside the bell. These pop up outside the browser, with the customer's name and a preview of what they said. **Click the notification to jump straight to that conversation.**

WHY A DESKTOP NOTIFICATION, WHEN THERE IS ALREADY A BADGE AND A CHIME

The badge only helps if you can see the tab, and the chime only helps if your sound is on and you are at your desk. A desktop notification is the one that reaches you when Pitstop is behind another window, or behind another application entirely. It is the alert for the tab you cannot see.

For that reason, notifications **do not appear while you are already looking at Pitstop** — the conversation list in front of you has already told you.

The first time you turn desktop notifications on, your browser will ask permission. You have to allow it: the toggle cannot work without it, and Pitstop only asks at the moment you click, never on its own.

IF THE BUTTON SHOWS RED, OR WILL NOT TURN ON

That means notifications are blocked for this site in your browser, which Pitstop cannot undo for you. Open your browser's site settings for Pitstop, set Notifications to Allow, then click the button again. If the button is greyed out entirely, your browser does not support them — Safari on iPhone and iPad, for example.

If you allowed notifications before and they later stopped, the toggle will show as off. That is intentional: it reflects what your browser will actually let through, not what you once chose.

Several messages from the **same** customer replace one another rather than stacking, so stepping away for ten minutes leaves you with one notification per conversation, not forty. Different customers each get their own.

Notifications follow whichever list you are viewing. On **Mine** you hear about your own conversations; on **All** you hear about everyone's. This matches how the badge and the chime already behave.

Keyboard

With the conversation list in focus, **↑** and **↓** move the highlight down the list and **Enter** opens the highlighted conversation. The highlight follows whatever you last clicked, so you can click into the list and then keep going from the keyboard. Typing in the reply box is never interrupted — the arrow keys go back to moving the cursor while you are writing.

Replying

Every channel is two-way: type in the composer and send. The toolbar has an **emoji** picker, **canned responses** (or type **/** in the box to search them), and, on Instagram and Messenger, a **paperclip** to attach an image or PDF. WhatsApp replies are text only for now. Switch to **Private note** to record an internal handover note (amber, never sent to the customer); notes support **bold** and *italic*. When you are finished, **Done** closes the chat; it **reopens** by itself if the customer writes again.

On an **email** conversation the composer adds the header fields you expect: **To** (pre-filled with the customer, editable) and a **Subject** (pre-filled from the thread, editable), plus a **+ Cc / Bcc** link that opens those two fields when you need to copy a colleague or vendor. Put more than one address in any field separated by commas. The reply sends from the care mailbox and stays on the same email thread.

Files a customer emails you appear as chips under their message; click one to open it. A chip that will not open always explains itself when you hover: it may be too large, or beyond the ten-file limit, or simply not fetched yet. **Attachments are kept in Pitstop for six months and then removed**, after which the chip reads "Kept for 6 months, then removed". Nothing is actually lost when that happens: every original is still on the email itself in Gmail, which is where to look for anything older.

① WHATSAPP: REPLY INSIDE THE 24-HOUR WINDOW

You can now reply to WhatsApp chats right here. The customer's messages are pulled in live, so you see both sides of the conversation, not just ours, and the full chat history scrolls back, not only the most recent messages. Free-text replies work while the 24-hour window is open (the header shows whether it is). Once that window closes the composer is disabled until the customer writes again; pre-approved templates to reopen a closed chat are coming next. Attachments on WhatsApp are not supported yet.

MESSAGE LENGTH: INSTAGRAM AND MESSENGER STOP AT 2,000 CHARACTERS

Instagram and Messenger refuse anything longer than 2,000 characters, and WhatsApp stops at 4,096. As you approach the limit a counter appears under the composer, in amber; go past it and the counter turns red and **Send** greys out until you shorten the message or split it in two. Private notes have no limit, because they never leave Pitstop. Before this counter existed the send simply failed with no explanation, so if you remember retrying a long reply several times, that is what was happening.

IF YOU SEE "SHOWING A MIRRORED COPY"

A yellow banner above the messages means Pitstop could not reach the live conversation and is showing its own stored copy, which may be missing the customer's most recent messages. **Do not assume the last message you can see is the last one they sent.** Press **Retry** in the banner; if it keeps failing, flag it rather than replying blind.

Linking and creating tickets

A conversation can be tied to a ticket so the full history lives together. If a ticket already exists, use **Link ticket** and enter its number. If there is none yet, use **Create ticket**: it makes a fresh ticket from the conversation, filling in the channel and the customer for you, and links it in one click. Once linked, the header shows the ticket number as a button straight through to it.

▲ THE 24-HOUR WINDOW

On every channel you can reply freely only within 24 hours of the customer's last message, and the header shows whether the window is open. On Instagram and Messenger a reply outside it still sends, tagged as a human-agent follow-up. On WhatsApp the composer is disabled outside the window (templates to reopen are coming), so reply promptly while it is open.

Post Comments

/comments

AGENT

LEAD

ADMIN

The Comments screen is where public comments left on our Instagram posts are read, answered and cleared. Unlike a direct message, a comment is visible to everyone who looks at the post, so an unanswered question sits in front of every future shopper who scrolls past it.

WHAT IT'S FOR	READ OR ACT?	WHO	LAYOUT
Answering and moderating public post comments.	An action screen: reply, hide, close.	AGENT LEAD ADMIN	Comment list · comment detail.

How comments reach this screen

Pitstop checks Instagram for new comments **every ten minutes** and lists them here. That is different from the Inbox, where a direct message arrives within seconds, and it is worth remembering: a quiet Comments screen may simply be a few minutes behind rather than genuinely empty. If you want the latest right now, press **Sync now** and the list refreshes from Instagram immediately.

Replies work the other way round: when you reply here, it posts to Instagram **straight away** and appears on the post publicly, under the L.O.T account.

① INSTAGRAM ONLY, FOR NOW

Facebook Page comments are not included yet. Reading them needs a permission our app has not been granted, so until that is approved, Facebook comments still have to be handled in Meta Business Suite.

Working the list

The controls along the top narrow the list: **Open / Done / All**, then **All / Mine / Unassigned**, then a search box that matches the comment text or the commenter's handle. The counter on the right shows how many comments are open and how many of those nobody has taken.

- 1 Pick a comment** from the list. The right-hand side shows it in context, with the post it was left on and any replies underneath.
- 2 Claim it** so the rest of the team can see it is yours. Replying claims it automatically, so this is only needed when you want to take something before you answer it.

- 3 **Reply** in the box at the bottom. Your reply is public and posts under the L.O.T account, not your own.
- 4 **Mark it Done** when it needs nothing further. Done is about your work, not about the post: a comment you have marked Done is still visible to the public exactly as before.

△ DONE AND HIDE ARE TWO DIFFERENT THINGS

Done clears the comment off your queue and changes nothing on Instagram. **Hide** actually removes the comment from public view on the post. A rude or abusive comment usually wants both; an ordinary question that you have answered wants only Done. Hide can be undone at any time with **Unhide**.

Replying inside a thread

Instagram keeps comment replies one level deep, so if you answer somebody who was themselves replying to another comment, your reply is added to the top of that group rather than nested underneath their message. This is Instagram's behaviour, not a limitation of Pitstop, and the person you are answering is still notified.

Deleting a comment

🗑 DELETE IS PERMANENT AND ADMINS ONLY

Deleting removes the comment from Instagram for good. It cannot be undone, by us or by Meta, and the person who wrote it may notice. Only admins see the button, and it asks for confirmation before acting. **Reach for Hide first:** it takes the comment out of public view just as effectively and can be reversed if the call turns out to be wrong.

What is not here yet

A few things are deliberately left for later so the screen stays simple: comments are not auto-assigned to anybody (you pick them up yourself), there is no way to raise a ticket directly from a comment, and there is no private-message reply to a commenter. If any of those would help you, say so and they can be added.

Calls

/calls

AGENT

LEAD

ADMIN

Calls is the phone-call log: who called, who answered, what was missed, and which calls still need a callback or a ticket. Agents use it to clear missed calls and turn them into tickets.

WHAT IT'S FOR	READ OR ACT?	WHO	UPDATES
The call log and missed-call follow-up.	Read; quick row actions.	AGENT LEAD ADMIN	List every 30s; new calls appear within about 2 minutes.

Reading it

Six cards show Calls today, Answered, Missed, **Abandoned**, Answer rate and Awaiting callback. Tabs split the list into All Calls, My Calls, Unassigned, **Open**, **Closed**, Missed, **Abandoned** and **Needs callback**. Open and Closed show calls whose linked ticket is still being worked or has been resolved; calls with no ticket yet stay under All, Unassigned or Missed. Each row shows the time, direction (incoming or outgoing), phone, customer, agent, talk time, a status badge (**Answered**, **Missed**, **Abandoned**, or Live), and the linked ticket if there is one.

▲ MISSED AND ABANDONED ARE NOT THE SAME THING

Missed means nobody picked up. **Abandoned** means the caller reached us and hung up before anyone spoke, usually within a few seconds. They need different responses: a missed call is a callback, an abandoned call usually means the person could not get through and may already have tried more than once. Both appear in **Needs callback**.

Older calls from before 20 August 2026 do not have this split. The previous phone system reported almost every abandoned call as answered, so historic Missed counts read far lower than reality. If you are comparing this month against last, expect the numbers to look worse; the earlier figure was the wrong one.

Talk time

The **Talk** column shows how long someone was actually speaking, with the total time the line was open shown smaller after a slash. So `1:45 /2:02` means a 2 minute 2 second call with 1 minute 45 seconds of conversation; the rest was ringing and connecting. A dash in the first position means nobody spoke at all.

Calling someone back

The softphone — calls ring inside Pitstop

Since 24 August 2026, agents with a SIP device take calls *in the browser*. A small bar sits at the bottom-left of every Pitstop page: a green dot and “Softphone ready” means you are reachable. The first time, the browser asks for microphone permission — press Allow, or the softphone cannot register.

An incoming call shows the number with **Accept** and reject buttons; during a call the bar carries a timer, **mute**, **hold** and **hang up**. To dial out, press **Dial** on the green bar, type the number and press Call.

⚠ WORK CALLS FROM PITSTOP, NOT THE EXOTEL TAB

Log out of the Exotel dashboard before using the softphone — both register the same device, and two registrations fight each other. Also keep Pitstop open in *one* tab: a second tab re-registers the device and calls may ring in the wrong one.

① IF THE BAR SAYS “SOFTPHONE OFFLINE”

Press **Retry**. If it stays offline, check the site’s microphone permission in your browser settings, then tell your lead with a screenshot.

Every row has a **Call** button. Press it and *your* phone rings first; answer it and we then dial the customer and join you together. The customer sees the Legend of Toys number, never your personal one, and you never have to type a number.

If nothing happens, you may not have a calling device set up yet. Ask an admin to add you on Admin then Telephony.

Acting on a call

① MARKING A CALL BACK CLEARS IT FROM THE LIST

Once you press **Mark Called Back**, the call leaves both the **Abandoned** and **Needs callback** tabs. That is how those lists stay as worklists rather than growing forever. The call itself is not gone — it is still under **All Calls**, with the time and the person who returned it recorded against it.

The **Abandoned** figure on the cards at the top does *not* go down, and that is deliberate: it counts how many people could not get through today, and a call you returned still happened.

There is no undo for this one yet, so press it only once you have actually made the call. If you mark one by mistake, find it under All Calls and tell your lead.

The row menu (...) offers the right actions for that call. A call whose ticket needs nothing can be closed in one press with **Nothing needed – close**; an **Undo** button appears next to the ticket number afterwards if you change your mind. A missed or abandoned call with no ticket can be **Marked Called**

Back or turned into a ticket with **Create Ticket**; a call with a ticket offers **Open Ticket**. Clicking a row opens its ticket if it has one, otherwise the call detail.

① **EVERY CALL GETS A TICKET, AND THAT IS ON PURPOSE**

Even a very short call opens a ticket. That is deliberate: someone who rang and hung up after four seconds usually could not get through, and the ticket is what a later WhatsApp message or callback attaches to. If it truly needed nothing, close it with one press rather than leaving it sitting in the queue.

The list shows 50 calls at a time. Use the **Prev** and **Next** buttons at the bottom to page back through older calls; the footer shows which rows and page you are on. Changing the tab, filters or search starts you back at the first page.

① **PHONE IS SHOWN IN FULL HERE**

Unlike the queue, the calls list shows the full phone number so you can call back. Turning a missed call into a ticket carries the call's details into the new ticket and links the two together.

① **REPEAT CALLS STAY ON ONE TICKET**

When the same customer calls again about something already open (a dropped call, a callback, a same-day follow-up), the new answered call attaches to their existing ticket instead of opening a duplicate, so you don't get a flood of near-identical tickets. Every call is still logged here as its own row, and the ticket's history notes that the repeat call was folded in.

Call Detail

</calls/detail>

AGENT

LEAD

ADMIN

Call Detail is the full record of one call: its timeline, the agent, and the recording, plus the buttons to turn it into a ticket or mark it called back.

WHAT IT'S FOR One call's full record.	READ OR ACT? Read; two actions.	WHO AGENT LEAD ADMIN	OPENED FROM The Calls list.
--	------------------------------------	----------------------------	--------------------------------

What it shows

A header card with the customer, status and the call's phone, duration, MyOperator account and department. A **timeline** card (created, answered, ended, called back, ticket created). An **agent** card. And, when one exists, a **recording** card with a link to play it.

Acting on it

Create Ticket From Call appears when no ticket is linked and opens a pre-filled New Ticket. **Mark Called Back** appears for a missed call that has not been returned. Once a ticket is linked, both buttons disappear and the ticket shows in the timeline.

Listening to the recording

Press **Load recording** and a player appears. Press play. The link is only good for about an hour, so if you come back to the page later, reload it and press Load recording again to get a fresh one.

The page also has a **Call back** button. Your phone rings first, then we dial the customer and join you together.

① SOME CALLS HAVE NO RECORDING

Very short and unanswered calls often have nothing recorded, and the page will say so. Calls from before 20 August 2026 cannot be played at all: the old phone system only ever sent us a file name, never the audio, so those show a note instead of a player. That is a limitation of the old system, not a fault here.

PART 3

Analyze

Cross-cutting reports on tickets and calls.

Support Analytics

/analytics

LEAD

ADMIN

Support Analytics is the “what are people complaining about” screen: it counts **complaints**, not every ticket, and breaks them down by product, issue, product line and channel over a date range. Open it to spot a pattern, not to reconcile totals against Reports.

WHAT IT’S FOR	READ OR ACT?	WHO	RANGE
Complaint analytics: what, on what product, from where.	Read; two CSV exports.	LEADADMIN	Today / MTD / Last month / This year / Custom.

COMPLAINTS ONLY, NOT EVERY TICKET

This screen counts tickets that have been given an **issue category**. A general question or a call with no category attached is not a complaint and is left out. That is deliberate, and it means the totals here will not match the Tickets tab on Reports, which counts every case. If a number here looks low next to Reports, this is why.

KPI band

Four figures at the top: **Total Complaints**, **Within 3 Days** and **After 3 Days** (of purchase), and **Ageing Unknown**. Ageing is the complaint date minus the purchase date; “unknown” means there is no purchase date on file for that ticket, so its age cannot be worked out. These four cards are headline totals only; the **Ageing Trend** panel further down is where ageing is tracked over time.

Filters, sort and the “X of Y” counter

Six filters sit under the date range: **Product**, **Issue category**, **Product line**, **Sale channel**, **Support channel** and **Agent** (which lists **Unassigned** as its own option). They combine — pick a product and a channel together and the screen narrows to complaints matching both. Each dropdown’s options come from the whole date range, not from what the other filters have already narrowed to, so choosing a product never empties the issue-category list.

As soon as any filter is on, a counter appears next to **Clear filters**: “X of Y complaints in range.” That way a filtered view is never mistaken for a quiet month — the second number is what the date range holds before the filters are applied.

A single **sort** control (Most first / Fewest first / A–Z) reorders every ranked panel on the page and the product table together, so the same order applies everywhere at once.

▲ "TOP ISSUE SUB-CATEGORIES" IS THE TOP 20 BY VOLUME

That panel only ever holds the twenty biggest sub-categories in the range — it says so on screen. Sorting it reorders those twenty; it does not reveal sub-categories outside the top 20, however the sort is set.

Panels

Below the KPI band: a **Complaints by Product** matrix (rows are products, columns are issue categories, darker cells are higher counts), then ranked panels for **By Issue Category**, **Top Issue Sub-categories**, **By Product Line**, **By Sale Channel** and **By Support Channel**, and finally four trend panels with their full table underneath: **Product Issue Trend**, **Complaint Rate**, **Ageing Trend** and **Category Trend**.

A **Trend: monthly / Trend: weekly / Trend: daily** dropdown sits above the trend panels. Weekly buckets are week-commencing **Monday**. Switching grain carries through everywhere: all four panel headings, each table's first column, the chart axes, and the CSV export.

Complaint Rate

Bars are the number of complaints raised in each bucket; the line is a trailing rolling average over the buckets immediately before it, so you can see whether a spike is a one-off or part of a rising trend. The averaging window is counted in **buckets**, not calendar days, and the legend names it in buckets for exactly that reason: **7** on the daily view means the last 7 buckets that have data, **4** on weekly and **3** on monthly. This matters because a bucket with no complaints at all is left off the chart rather than plotted as zero.

Ageing Trend

Tracks the same three ageing groups as the KPI cards (within 3 days of purchase, after 3 days, unknown) over time, which the KPI cards alone do not do. The chart plots each group as a **share** of that bucket's complaints, not a raw count — raw counts track total volume, so a busy month would look like a worsening one even if the mix of early versus late complaints had not changed. The table underneath shows both figures for each bucket: the count and the percentage.

Export CSV and Export complaints

Two export buttons sit side by side, and they are not the same file.

Export CSV downloads exactly what is on screen: every panel, in one file — the KPI cards, the matrix, the ranked panels and the trend tables. It is aggregates only; it has never carried an individual ticket. The top of the file records which filters were active, how the panels were sorted, the complaint count for the filtered cohort and the count for the whole date range before filters — so a spreadsheet opened next week is never ambiguous about which slice of complaints it holds.

Export complaints downloads one row per complaint in the current cohort: the same date range and the same dimension filters applied on screen, saved as `pitstop-complaints-<from>-to-<to>.csv`. This is the file to reach for when a number on the dashboard needs explaining, because it carries the ticket, raised and closed timestamps, customer and order details, product and category fields, sale and support channel, agent, stage and disposition, the closing reason, and — unlike Export CSV — the

complaint description itself, the closing note, the inspection note and every internal note on the ticket (all of them in one cell, each stamped with its time and author, separated by " | "). Phone numbers are masked with the last three digits hidden, as in the Queue export; the customer's email address and order ID are included in full, which the Queue export does not carry. It is an internal file — treat it as one. Its header block names the range, every applied filter and the two cohort counts, exactly like Export CSV. The button reads **Preparing...** while it fetches, since it is a heavier download than the dashboard CSV.

▲ A VERY WIDE RANGE CAN BE MARKED INCOMPLETE

If a date range holds more than 50,000 complaints, the screen shows a warning that the figures below and any export are incomplete, and asks you to narrow the dates. Both CSVs carry the same warning if you export anyway. This cannot happen at today's volumes, but if you ever see the banner, that is what it means.

This screen needs the reports permission (the same one Reports uses), so it is visible to Lead and Admin only.

Reports

/reports

LEAD

ADMIN

Reports is the “how are we doing” screen: how many cases and calls happened over a date range, what they cost, and how products, platforms, agents and departments compare. Open it to review performance.

WHAT IT’S FOR	READ OR ACT?	WHO	RANGE
Ticket and call analytics.	Read; CSV export.	LEAD ADMIN	From / To (defaults year-to-date).

Three tabs

Tickets shows total cases and the three cost lines (return logistics, replacement units, refund payouts), a monthly trend chart coloured by outcome, and a breakdown by product, platform and agent. The agent table’s first column is **Raised**, not the old “Total”. **Calls** shows total inbound calls, **Reached an agent** and **Didn’t reach an agent**, an **Inbound answer rate**, and average duration, with breakdowns by department, account, direction and agent, plus an hourly distribution. **Agents** is the conversation report: a **Queries** KPI card sits above a per-agent table of **Assigned**, **Handled** and **Open**, plus replies, closures and response times across every channel.

On the Agents tab, hover any column heading (dotted underline) or summary card for its exact definition — the wording was agreed with the CS team, so what you read is what the number means.

Agent, Channel and Tag filters

The Agents tab has three combinable filters: **Agent**, **Channel** and **Tag**. All three are multi-select, so you can pick several agents at once — “these three agents” is the normal question on a tab that is entirely about agents, and one-at-a-time could not answer it. Leaving a filter with no chips selected means **All** on that dimension; it is not the same as picking every option one by one, and it is what the tab shows before you touch anything.

△ THE KPI CARDS AND THE WAIT BREAKDOWN FOLLOW THE FILTER TOO

Filtering the Agents tab does not just narrow the table beneath it. The KPI cards at the top (Queries and the rest) and the **What the resolution time is made of** panel lower down both recompute over the same filtered cohort. Filter to two agents and the headline numbers are those two agents' numbers, not whole-team totals sitting above a filtered table — the opposite is a common and believable thing to assume, so it is worth stating plainly.

One number behaves in a way worth knowing: **Handled** counts conversations an agent replied on, and under a filter it counts only the ones *assigned* to her. So her Handled reads **LOWER** filtered than on the unfiltered screen — the difference is conversations she replied on that belong to someone else. On 2026-09-04 that gap ran from about **3% to 11%** across the team. ⚠ For an agent who has replies but *no* conversations assigned to her, filtering to her shows **zero** Handled, not a small number — her work is real, it simply all sits on other people's conversations. Both figures are correct; they answer different questions. Her row does *not* change when you add a colleague to the filter.

A **Clear filters** button appears once any of Agent, Channel or Tag has a selection, and clears all three at once. The **Business hours only** checkbox sits alongside them and is unchanged: it is **on by default**, and response and resolution clocks run only during each agent's own rostered shift while it is ticked. Untick it for the around-the-clock view — the same figures with the clock never stopping — which is the customer's experience of the wait rather than the agent's. The basis in use is named under each time card and in the CSV export.

The **Export CSV** button on this tab records the Agent, Channel and Tag cohort at the top of the file (each as "All" or the selected names), alongside the Business hours basis, so a file opened a week later is never ambiguous about which slice it holds.

△ RESOLUTION CHANGED MEANING ON 24 AUGUST 2026

Resolution is now *Resolved ÷ Queries*: conversations genuinely resolved, as a share of queries raised. It previously counted *every* closure — including spam and duplicates — which overstated the rate by roughly three times. If you compare against a figure from before this date, expect a sharp drop; the earlier number was the inflated one. The old closure-based series still appears in the CSV export as "Closed rate %".

△ QUERIES COLUMN REMOVED FROM THE AGENTS TABLE

The per-agent table on the Agents tab no longer has a **Queries** column — **Assigned**, **Handled** and **Open** remain. The **Queries** KPI card above the table stays: Resolution rate and Resolve rate still divide by it.

RAISED CAN BE LOWER THAN CLOSED, AND THAT IS FINE

On the Tickets tab, the agent table's **Raised** column counts tickets raised in the range and assigned to that agent; **Closed** counts tickets they closed in the range, whenever those were originally raised. An agent clearing older backlog closes more than was raised to them in the same window — that is not a fault, it is how the two columns are defined.

▲ "MISSED" IS NOW "DIDN'T REACH AN AGENT"

The Calls tab's **Answered** and **Missed** cards are now **Reached an agent** and **Didn't reach an agent** (both inbound), alongside a new **Inbound answer rate**. The old "Missed" figure was wrong for the MyOperator era: that system recorded a call as missed in one rare case only, so it showed 45 missed calls for the whole period while 4,330 inbound calls never reached anybody. **Didn't reach an agent** also includes people who hung up in the phone menu before routing, so it is not all calls the team failed to answer.

Set the From and To dates at the top; changing either refetches immediately. A **Last 7 days** button sits next to the pickers, on every tab — it sets the range to today and the six days before it, the range the CS lead opens every morning. It is a shortcut, not a default: the Tickets tab still opens on year-to-date until you click it.

Daily trend (Agents tab)

Below the per-agent table on the **Agents** tab, a **Daily trend** panel charts one metric per day: a line for **All agents** and one line per agent, for whichever metric you pick from the dropdown — Queries, Answered, Never answered, Avg first reply, Avg resolution, Resolve rate, Handled, Resolved or Assigned. It follows the same Agent, Channel and Tag filters and the same Business hours checkbox as the KPIs above it, because each point on the chart is the Agents report computed for that single day. So every number keeps the date basis the rest of this tab already documents: Queries, Answered, Never answered and the averages/rates count conversations **raised** that day; Handled counts the day the agent replied; Resolved counts the day it was closed; Assigned counts the day it was handed over. One consequence worth knowing: a conversation replied to on three different days is Handled on each of them, so the daily Handled figures add up to more than the Handled total for the range (checked on 4 Sep 2026: 1,589 across seven days against 1,179 for the week), whereas Queries, Answered, Assigned and Resolved add up exactly. The panel repeats this note whenever Handled is selected.

Only the **six busiest agents** (by their Handled figures added up across the days shown — not the range total, which is the different number the note above explains) get their own line. For a count metric (Queries, Answered, Handled, and so on) everyone else is summed into an **Others** line, so the total still adds up. For a minutes or rate metric (Avg first reply, Avg resolution, Resolve rate) the smaller agents are simply **not drawn** — the panel says so in its caption — because an average of averages is not a number anyone should read.

UP TO 62 DAYS AT DAILY GRAIN

The trend plots one point per day, so a very wide range asks for a lot of points. When the selected range is longer than **62 days** — the tab opens on year-to-date, for instance — the panel plots the **last 62 days of that range** and says so in its caption. The KPIs and the table above it keep the full range.

The Agents tab's **Export CSV** now ends with a **Daily trend** block: one row per day for **All agents**, then one row per day per agent, with every metric as its own column (minutes and percentages labelled in the header). It is appended only when the panel has actually loaded — a failed or still-loading trend simply leaves the block off the file rather than exporting blanks that would read as zeros.

WEEKLY, MONTHLY AND CALLS TRENDS ARE NOT BUILT YET

The Daily trend panel is daily grain only, and only on the Agents tab. A week/month rollup and an equivalent trend for the Calls tab are not available.

What the resolution time is made of

On the **Agents** tab, the panel **What the resolution time is made of** breaks "Avg to close" into the parts it is actually built from. It answers the common objection that the figure is unfair because it counts time nobody was working.

FIGURE	WHAT IT COUNTS
Total to close	The full wall clock from the conversation being raised to it being closed. This has not changed — it is the same number as "Avg to close" in the table below.
– Waiting on customer	Every stretch where we had replied last and the next move was theirs.
– Waiting to be closed	Time after the <i>final</i> message in the conversation — how long it sat finished before someone pressed Close.
= Active handling	What is left: the part that reflects how fast the conversation was actually worked.

The same two figures appear per agent in the table as — **on customer** and — **unclosed**, and in the CSV export.

THE CLOCK IS NEVER PAUSED – THAT IS DELIBERATE

Waiting on the customer is *not* deducted from the headline. The real total stays the real total, and the waits are shown beside it so you can subtract them yourself when that is the fairer comparison. Quote whichever you mean, and say which one it is.

▲ CLOSING DISCIPLINE IS USUALLY THE BIGGER HALF

Measured across **7,118 conversations** in the settled 30-day window (24x7 basis, measured 2026-09-03): of **51.7 hours** to close, **12.9 hours (25%)** was waiting on the customer, **22.1 hours (43%)** was waiting to be closed, and **16.7 hours (32%)** was active handling. So a long “time to close” is still usually not customers being slow — it is conversations finishing and staying open, which remains the largest single component and still runs well above the customer-wait share. Closing promptly is what moves this number.

An earlier reading of this cohort (creation-dated, conversations closed in the week to 25 August 2026, showing 66%) has been superseded: it was taken while that cohort was still maturing and understated waiting-on-customer while overstating waiting-to-be-closed. Use the 30-day figures above.

IF THE PANEL IS MISSING

It is calculated separately from the rest of the page, so on a very wide date range it can appear a moment after the tables. If it does not appear at all, the figures did not reconcile with the main report and it hides itself rather than show numbers that disagree — the rest of the report is unaffected.

▲ EXPORT IS TICKETS ONLY

The **Export CSV** button exports the *ticket* figures (product, platform, agent and a cost summary) from the Tickets and Calls tabs — there is no separate calls CSV. The **Agents** tab is the exception: there the button exports the agent conversation report, including both “Closed rate %” and “Resolution rate %”. In the agent call table, “missed to returned” counts missed calls that were later called back, not still-missed ones. This screen needs the reports permission.

PART 4

Admin

Departments, MyOperator call accounts, and WhatsApp templates.

Departments

/admin/departments

ADMIN

Departments define the CS teams (Inbound, Outbound, Confirmation, Messaging and so on) and which agents belong to each. Departments control visibility: a non-admin only sees the tickets and calls in the teams they belong to. This screen is also where you set a teammate’s CS role.

WHAT IT’S FOR	READ OR ACT?	WHO	EFFECT
Teams, who is in them, and CS roles.	Action: create, toggle, assign.	LEAD ADMIN	Scopes everyone’s access.

Departments

The **Departments** table lists each team with its slug, name, sort order, agent count and an Active toggle. **Add Department** creates one (a URL-safe slug, a name, a sort order).

Agents: role and departments

The **Agents** table lists everyone. For each person you can set two things:

- 1

Set Role. Pick **Viewer**, **Agent** or **Team Lead**. A Viewer can open the queue but not work tickets; an Agent works tickets and calls; a Team Lead can also assign across the team, approve refunds and replacements, and see reports.
- 2

Departments. Tick every team the person belongs to. A person can be in more than one team and will then see the tickets and calls of all of them. Untick every box to remove them from all team views.

① YOU CAN ONLY SET CS ROLES HERE

The role options are limited to the CS tiers (Viewer, Agent, Team Lead). Anyone who holds a wider company role, such as a Garage admin, shows as **Garage-managed** with no dropdown and is changed in Garage, not here. You also cannot change your own role, to avoid locking yourself out.

⚠ **CHANGES APPLY INSTANTLY, NO CONFIRM**

Toggling a department active, changing a role, and ticking or unticking departments all take effect the moment you click, with no confirmation, so a misclick is live immediately. Removing someone from every department drops their tickets and calls out of all team views. A slug is fixed once the department is created. The person may need to refresh the page for a new role or team to take effect.

MyOperator Accounts

/admin/myop

ADMIN

This page has two parts. The **Exotel** panel at the top shows whether the live phone connection is healthy and lets an admin recover calls Pitstop missed. Below it, MyOperator Accounts registers each calling account that feeds phone calls into Pitstop. Each account gets its own webhook URL and secret; you set it up here, then point MyOperator at the generated URL so its call events flow in.

WHAT IT'S FOR	READ OR ACT?	WHO	WHEN
Connecting calling lines to Pitstop.	Action: create + activate.	ADMIN	Onboarding a new number.

Adding an account

Add Account captures a slug, a name, and optionally the DID (caller number) and owner email. Each account row shows its generated **webhook URL** and **env-var name**, each with a copy button, and an Active toggle. There is no edit or delete: an account can only be created and then activated or deactivated.

ⓘ

CREATING THE ACCOUNT IS ONLY HALF THE SETUP

Calls will not authenticate until an engineer sets the matching secret in Cloudflare (the `MYOP_WEBHOOK_SECRET_<SLUG>` shown on the row), using `wrangler secret put`. That is a one-time technical step outside Pitstop. Because there is no edit, a typo in the slug, DID or owner can only be fixed by creating a new account.

The Exotel panel

A green dot and **Connected** means Pitstop is talking to the phone system and shows when the last call arrived. Red means it is not, and the message says why; the fix is usually a credential, so send that message to whoever set the account up rather than retrying.

Run backfill recovers calls that never reached Pitstop, for example during an outage. Set the date and time to start from and press it. It writes call records only and creates **no tickets**, so recovering a day of history will not flood the queue or reset anyone's clock. Running it twice is safe: it never duplicates a call.

① MYOPERATOR IS KEPT, NOT DELETED

Phone calls moved from MyOperator to Exotel on 19 August 2026. The MyOperator accounts below are left in place on purpose, so that if any number turns out not to have moved, its calls still reach Pitstop. Leave them active.

Telephony

/admin/telephony

ADMIN

Telephony decides who can place calls from inside Pitstop, and which phone rings when they do. Someone who is not listed here cannot use the Call button anywhere in the app.

WHAT IT'S FOR	READ OR ACT?	WHO	WHEN
Letting people call customers from Pitstop.	Action: set up and edit.	ADMIN	Someone joins, leaves or changes number.

Reading it

The line at the top tells you how many people can currently place calls, and warns you in orange if any cannot. The table lists everyone who might need to make calls: agents, leads and admins, plus anyone already set up. People who cannot call are listed **first**, because they are the ones needing attention.

Each row shows the person, their role, their SIP device, their mobile number, and which of the two rings first.

Setting someone up

Find the person in the list. If they have never been set up, their Rings column reads *cannot call* and the button says **Set up**.

Press Set up (or **Edit** for someone already configured). The row becomes editable.

Fill in what they have. The mobile number is enough on its own. The SIP device is optional and only exists for people set up for browser calling; you can find it in Exotel under Co-workers and Groups.

Choose what rings first: **Mobile** calls their phone, **Softphone** calls their computer.

Press Save. They can use the Call button straight away; there is nothing to restart.

The softphone setup button

Above the roster sits **Run softphone setup**. It registers every active SIP agent with the calling platform so the in-app softphone works for them. It is safe to press again at any time — existing agents are skipped — and the per-person result is shown underneath, so a rejected name or a missing email is visible immediately rather than silently dropped. Press it after adding a new SIP agent to the roster.

How a call actually goes out

When someone presses Call, we ring *them* first. Once they answer, we dial the customer and join the two together. The customer sees the Legend of Toys number, so nobody's personal number is ever exposed, and the call is recorded like any other.

① A MISSING SIP DEVICE IS NOT A PROBLEM

If someone is set to Softphone but their computer is not registered, the call falls back to their mobile rather than failing. So a person with a mobile number always gets through.

⚠ ENTER THE MOBILE NUMBER CAREFULLY

The number here is how we recognise that person on a call. If it is wrong, their calls will show as unassigned in the call log and their name will not appear on their own tickets. Enter it the way you would dial it; the page tidies up the format for you.

① TWO PEOPLE CANNOT SHARE ONE SIP DEVICE

Each SIP device belongs to exactly one person. If you try to give the same one to a second person, the save is refused. Copy the device id from Exotel for the right person rather than reusing one that looks similar.

WhatsApp Templates

/admin/wa-templates

ADMIN

WhatsApp Templates manages the library of pre-approved messages agents use to start a conversation or message a customer outside the 24-hour reply window. Open it to add or edit a template’s wording, label, category and language.

WHAT IT’S FOR	READ OR ACT?	WHO	CATEGORIES
The approved-message library.	Action: create / edit / toggle.	ADMIN	Utility · Marketing · Authentication.

Managing templates

Each row shows the template **name** (the Meta-approved name), the **label** agents see, its category, language, how many **placeholders** (`{{1}}` , `{{2}}` ...) the body has, and an Active toggle. **Add Template** or **Edit** opens a form for the name, display label, category, language and body; the placeholder count is counted from the body automatically as you type.

⊖ THE NAME MUST MATCH META EXACTLY

A template’s **name** must match the name approved in Meta Business Manager, or it will not send; and the name is locked after creation (fix a wrong one by making a new template). Today this screen *records* templates and composes wording; the actual WhatsApp send is a later phase, so saving a template here does not yet deliver a message on its own.

Meta Diagnostics

/admin/meta

ADMIN

Meta Diagnostics answers one question: **why did Instagram or Messenger refuse to send a reply?** When an agent’s message to a customer bounces, the refusal is recorded here with Meta’s own error, and you can probe the chat to see whether the problem is our connection or that particular customer.

WHAT IT’S FOR	READ OR ACT?	WHO	RECORDS FROM
Diagnosing refused Instagram / Messenger replies.	Read only. It never sends a message.	ADMIN	26 Aug 2026 onward.

Reading the list

Each row is **one chat and one error code**, not one message — if an agent pressed send four times and it failed four times, that is a single row with **Attempts 4**. **Window** says whether the 24-hour reply window was open *at the moment of the attempt*, which is the fact that separates a normal out-of-hours refusal from a genuine fault. Use the range selector to look further back.

What the error codes mean

CODE	MEANING	WHAT TO DO
10	Outside the 24-hour window and our “Human Agent” permission is still pending Meta review.	Nothing. This is expected behaviour, not a fault. The chat reopens if the customer writes again.
100	Invalid parameter — almost always a reply longer than 2,000 characters.	Shorten the reply. The composer now warns before this happens.
190	Our access token has expired or been revoked. Affects <i>every</i> chat, not one.	Tell Afshaan. Replying from the Instagram app still works meanwhile.
200	Permissions error, refused for that one recipient while other chats send normally.	Press Diagnose on the row, then read the warning under Diagnose — as of 27 Aug 2026 the probe rules the token out but cannot yet name this cause. Pass the result to Afshaan.

Diagnose

Diagnose makes three read-only checks against Meta using the same credential the reply would have used: one for our own account (`me`), one for the customer (`recipient`) and one for the conversation itself (`conversation`). Read the result this way:

1. `me` fails — the problem is our token, not the customer. That is an Afshaan job, and it will be affecting everybody.
2. `me` succeeds but `recipient` or `conversation` fails — our connection is healthy and Meta is refusing this specific person. The response body is the evidence of why.
3. All three succeed — this rules the token out, and rules nothing in. See the warning below.

⚠️ THREE GREEN RESULTS DO NOT MEAN THE CHAT IS HEALTHY

On 27 August 2026 the probe was run on a chat that had just been refused four times with a code 200, and on two chats that were sending normally minutes earlier. **All three came back identical** — same three 200s, same credential, differing only in the customer's follower count and display name.

These checks prove the account *exists* and the conversation is *reachable*. They cannot see whether Meta will let us *message* it, which is the thing a code 200 is refusing. So report a clean probe as “the token is fine and the cause is still unknown” — never as “no problem found”.

① IT SENDS NOTHING

Diagnose is safe to press at any time, on any chat, as often as you like. It reads three things from Meta and writes nothing — the customer sees no message and the conversation is not touched.

⚠️ DO NOT TRANSLATE THE RAW RESULT INTO AN EXPLANATION

The probe output is shown exactly as Meta returned it, on purpose. In August 2026 a Meta error was explained from the documentation rather than from a measurement, and the wording that reached agents told them a reply would go through while it was in fact failing silently. If the result does not obviously say what is wrong, that is the honest answer — pass it to Afshaan rather than guessing at a cause.

Probing a chat by hand

Refusals are only recorded from 26 August 2026 onward, so an older report will not appear in the list. The box at the bottom takes a conversation id directly and probes it the same way.