

Relay Operations Manual

The self-serve guide to every screen in Relay, the system that talks to customers.

- VIEWER
- AUTHOR
- MANAGER
- APPROVER
- ADMIN
- SUPER-ADMIN

Version	Updated	Application
1.9.3	2026-09-11	relay.legendoftoys.com

Contents

GETTING STARTED

Welcome to Relay	5
Signing In & Getting Around	7
Roles & What You Can Do	9
The Send Gate	11

OVERVIEW

Overview /	15
------------------	----

SEND

Campaigns /campaigns	17
Audience Exclusions /campaigns	20
A/B Testing /campaigns	22
Experiment Log /experiments	24
Journeys /journeys	25
Bot Builder /journeys	28
Link Tracking & Attribution /campaigns	33

AUDIENCE

Activity /activity	36
Segments /segments	37
Contacts /contacts	40
Forms /forms	42

BUILD & MEASURE

Templates /templates	46
SMS & RCS /templates	48
Library /library	51
Links /links	52

Analytics /analytics 54

ADMIN

Users /admin/users 57

Roles /admin/roles 58

Approval & Caps /admin/settings 60

Sender Identities /admin/senders 62

Connectors /admin/connectors 63

Suppressions /suppressions 64

PART 1

Getting Started

What Relay is, how to sign in and get around, and what your role lets you do.

Welcome to Relay

- VIEWER
- AUTHOR
- MANAGER
- APPROVER
- ADMIN
- SUPER-ADMIN

Relay is the system that talks to Legend of Toys customers. Every marketing email, every WhatsApp message, every SMS that goes out to a customer leaves through Relay. This manual walks you through every screen, in the order you are likely to need them.

WHAT IT'S FOR	CHANNELS	WHO USES IT	WHERE
Sending the right message to the right customers, on purpose.	Email, WhatsApp, SMS.	The brand and CX teams.	relay.legendoftoys.com

The idea in one page

Relay holds a record for every customer we know about, called a **profile**. A profile gathers up everything we have learned: their email address and phone number, whether they have agreed to hear from us, what they have bought, and what they have done on the website.

You then do three things with those profiles:

- 1 **Group them.** A **segment** is a rule that picks out a set of profiles, for example "bought exactly once and has not been back in 90 days".
- 2 **Write something to send them.** A **template** is the message itself, one per channel.
- 3 **Send it.** Either as a **campaign** (a one-off broadcast to a segment) or a **journey** (an automation that reacts to what a customer does, like abandoning a cart).

Campaign or journey?

	CAMPAIGN	JOURNEY
Triggered by	You, pressing send.	Something the customer does.
Goes to	Everyone in a segment, once.	One person at a time, as they qualify.
Good for	A launch, a sale, a newsletter.	Cart recovery, order updates, win-backs.
Lives on	Campaigns	Journeys

Nothing sends without passing the gate

Every single message, no matter which screen it came from, is checked one last time before it goes out. That check is called the **send gate**, and it is the reason Relay is safe to use. It blocks messages to people who have opted out, messages that would arrive in the middle of the night, and messages to someone we have already contacted too often today.

The send gate has its own chapter, and it is worth reading before you send anything. It explains most of the times you will wonder “why did fewer people get this than I expected?”

IMPORTANT

Relay sends to real customers. There is no practice mode once test mode is switched off. Read **The Send Gate** and **Campaigns** before your first send, and use the **Send test** button to see the message on your own phone first.

Signing In & Getting Around

VIEWER

AUTHOR

MANAGER

APPROVER

ADMIN

SUPER-ADMIN

How to get in, what the sidebar groups mean, and the two shortcuts worth learning on day one.

SIGN IN WITH Your <code>@legendoftoys.com</code> Google account.	GROUPS Overview, Send, Audience, Build & measure, Admin.	SEARCH ⌘ K	IF A PAGE IS MISSING Your role does not have it. See Roles.
--	--	---	---

Signing in

- 1 Go to `relay.legendoftoys.com`. You will land on a sign-in screen.
- 2 Sign in with Google. Only `@legendoftoys.com` accounts are accepted. A personal Gmail will be refused.
- 3 If you see “Relay access required”. Your account is valid but has not been given a Relay role yet. Ask an admin to add you on the Users screen.

The sidebar

Overview	The landing screen. What is sending right now, what is waiting for approval, and recent activity.
Send	Campaigns, the Experiment log, and Journeys. Everything that puts a message in front of a customer.
Audience	Activity, Segments and Contacts. Who we know and what they have done.
Build & measure	Templates, Library, Links and Analytics. The parts a send is made from, and the results afterwards.
Admin	Users, Roles, Approval & Caps, Sender Identities and Connectors. Most people will not see this group at all.

Two shortcuts

⌘ K opens the command palette. Start typing a screen name or a campaign name and press Enter. On a **phone** there is no keyboard for that shortcut and the side rail is hidden, so tap the **magnifying-glass icon in the bar across the top** to open the same palette. It is the fastest way around the app.

Esc closes any open dialog without saving.

The bar at the top

The top bar shows where you are (group, then screen) and carries the waffle menu on the right, which jumps to the other Legend of Toys systems: Garage, Redline, Depot, Snorkel, Ignition, Docket, Pitstop, Odo and Throttle. Relay itself is pinned at the top of that menu as a reminder of where you already are.

✓ TIP

If a screen looks stale, refresh the page properly with **⌘ Shift R**. A normal refresh can serve you an older cached copy of the app.

Roles & What You Can Do

VIEWER AUTHOR MANAGER APPROVER ADMIN SUPER-ADMIN

Relay has its own list of roles, separate from every other Legend of Toys system. Being an admin in Garage or Pitstop gives you nothing here. This chapter explains what each role can do and why the split exists.

SIX ROLES Viewer, Author, Manager, Approver, Admin, Super-admin.	SET BY An admin, on the Users screen.	SEPARATE FROM Every other LOT system.	NO ROLE? You cannot open Relay at all.
---	--	--	---

The roles

ROLE	CAN DO
VIEWER	Open and read every screen. Changes nothing. Good for anyone who needs the numbers but should not touch the sending.
AUTHOR	Build segments, write templates, and prepare campaigns. Cannot press send. This is the safe default for a marketer learning the system.
MANAGER	Everything an Author can do, plus pressing send, scheduling a campaign, and activating a journey.
APPROVER	Signs off campaigns that are large enough to need approval. Can read everything, but does not build.
ADMIN	Everything above, plus adding people, managing sender identities and connectors, and editing a customer's consent.
SUPER-ADMIN	Everything, plus the controls that apply to the whole system: roles, the approval threshold, the frequency cap, quiet hours and the global send lock.

Why building and sending are separate

The split between **Author** and **Manager** is deliberate and worth keeping. Preparing a campaign is reversible: you can edit a draft as many times as you like. Sending is not. Once a broadcast starts fanning out to twenty thousand people, there is no recall.

So the person who writes the message and the person who releases it can be two different people, and for anything large they usually should be.

If a screen is missing

Screens you do not have permission for are hidden, not greyed out. If someone describes a screen you cannot find, you almost certainly do not have the role for it. The Admin group in particular is invisible to most people.

NOTE

Access can be switched off without removing you. If you were using Relay yesterday and today you get “Relay access required”, your access has been set to inactive rather than deleted. An admin can turn it back on in one click.

The Send Gate

VIEWER

AUTHOR

MANAGER

APPROVER

ADMIN

SUPER-ADMIN

Every message Relay sends is checked one last time immediately before it goes out. This chapter explains those checks in order, because they are the reason a send sometimes reaches fewer people than the screen predicted.

APPLIES TO	RUNS	RESULT IF BLOCKED	SET BY
Every message, every channel, every screen.	Per recipient, at the moment of sending.	The message is logged as <code>skipped</code> , with a reason.	SUPER-ADMIN on Approval & Caps.

The checks, in order

0. Test mode	A global lock. While it is on, nothing reaches anyone outside the internal test list, no matter what you do. It is the master safety switch.
1. Suppression	A hard block on an address or number, for example after a permanent email bounce or an SMS opt-out. This one outranks everything, including order confirmations.
2. Consent	Marketing needs an explicit opt-in on that channel. No opt-in, no marketing message. Order updates are not affected.
3. Frequency cap	A limit on how many marketing messages one person can receive in a rolling window. Currently three in twenty-four hours.
4. Quiet hours	No marketing during the quiet window. Set per channel, because the rules differ: promotional SMS in India is legally restricted to 10:00 to 21:00, WhatsApp is not, and email is exempt entirely.
5. Channel rules	Basic sanity: a valid email address, a plausible phone number, and for WhatsApp, an approved template if the 24 hour reply window has closed.
6. Send budget	A daily ceiling on marketing sends, used to warm up a new sending number gently. Checked last, so a slot is never wasted on a message another check would have stopped.

What bypasses what

PURPOSE	CONSENT, CAP, QUIET HOURS, BUDGET	SUPPRESSION
marketing	All apply	Applies
service (a survey after the customer contacted us)	Quiet hours apply, the rest do not	Applies
transactional / utility (order confirmed, out for delivery)	None apply	Applies

The logic is simple: a customer who asked us to stop marketing at them still needs to know their parcel is arriving. But a customer whose email address hard-bounced cannot receive anything at all, because there is nowhere to send it.

Reading a skipped message

When the gate stops something, it is recorded rather than silently dropped, and the reason is shown in Analytics and on the contact. The common ones:

REASON	WHAT HAPPENED
no_consent	Never opted in to marketing on that channel.
freq_cap	Already had the maximum number of marketing messages today.
quiet_hours	It was outside the sending window for that channel.
suppressed	Hard block on the address or number.
budget_exhausted	The daily marketing send ceiling was reached.
test_mode_blocked	Test mode is on and this person is not on the test list.

⚠ WATCH OUT

A skipped message means the person got **nothing**. It is not a delivery failure and it is not a retry queue. If a campaign reports a lot of `budget_exhausted`, those people were never contacted and will not be contacted later unless you send again.

① **NOTE**

The audience number on the campaign screen accounts for suppression, consent and your own exclusion rules. It cannot account for the frequency cap or quiet hours, because those depend on the exact moment each message is sent. Expect the delivered count to be a little lower than the predicted audience.

PART 2

Overview

The screen you land on, and what its numbers mean.

Overview

/

VIEWER

AUTHOR

MANAGER

APPROVER

ADMIN

SUPER-ADMIN

The screen you land on. It answers one question fast: is anything happening right now that needs me?

WHAT IT'S FOR	READ OR ACT?	WHO	UPDATES
A live status board for the whole system.	Read-only. Everything on it is a link to somewhere else.	Everyone with Relay access.	Refreshes itself while a campaign is sending.

The status strip

Test mode	Reads LOCKED or OPEN . Locked means nothing can reach a real customer. This is the single most important thing on the screen: check it before you assume a send went out.
Sending now	Campaigns currently fanning out, with progress. A broadcast to a large audience takes hours, so seeing something here for a while is normal.
Awaiting approval	Campaigns submitted by an Author and waiting for an Approver. If this is not empty, someone is blocked.
In queue / In flight	Messages accepted and on their way to the provider. A small number here at any time is healthy.
Revenue	Money attributed to Relay sends. See Analytics for how attribution is worked out.

Recent activity

A running list of campaign and journey lifecycle events: submitted, approved, started, finished. It is the quickest way to see what a colleague did this morning without asking them.

✓ TIP

If you are ever unsure whether a send actually went to customers, look at Test mode here first. "It said it sent but nobody got it" is almost always test mode being locked.

PART 3

Send

The three ways a message leaves Relay: a one-off broadcast, an automated journey, and the log that tells you which worked.

Campaigns

/campaigns

VIEWER

AUTHOR

MANAGER

APPROVER

A campaign is a one-off broadcast: you pick an audience, pick a message, and send it. This is the busiest screen in Relay and the one where mistakes are most expensive, so it has the most guard rails.

WHAT IT'S FOR	READ OR ACT?	WHO	UNDO?
Sending one message to one segment, once.	Action screen. Sends real messages.	AUTHOR builds, MANAGER sends, APPROVER signs off.	No. A started broadcast can be cancelled, not recalled.

The lifecycle

A campaign moves through fixed states, and what you can change depends on where it is.

STATUS	MEANING	EDITABLE?
draft	Being built.	Yes, fully.
pending_approval	Submitted, waiting for an Approver.	No.
approved	Cleared to send.	No.
scheduled	Will send by itself at a set time.	No, but the schedule can be cancelled.
sending	Fanning out now.	No.
sent	Finished.	No.

⚠ WATCH OUT

A campaign becomes read-only the moment you submit it. That is deliberate: an approver must sign off the thing that actually sends, not a version that changed afterwards. If you need a change, have it rejected back to draft.

Building one

- 1
- Name it.** Use something you will recognise in Analytics in three months. “Diwali 2026 catalogue”, not “test 4”.

- 2 **Pick the channel and purpose.** Purpose matters: `marketing` is gated on consent and quiet hours, `transactional` is not. Do not label a promotion as transactional to dodge the gate.
- 3 **Pick the audience.** Choose a segment. The line underneath immediately tells you how many people will actually receive it.
- 4 **Set exclusions if you need them.** See the Audience Exclusions chapter.
- 5 **Pick the template.** Only templates for the chosen channel appear. On WhatsApp, only templates approved by Meta can be used; on SMS and RCS, only vendor-approved ones will actually send.
- 6 **Fill the campaign variables if the template uses them.** A campaign has no order or browsing event behind it, so values like a sale name, discount or code come from the campaign's variables box — e.g. `{"sale_name": "Freedom to Play Sale", "code": "FREEDOM5"}`. A missing value means that person's message is skipped, not sent with a hole.
- 7 **Send a test.** The **Send test** button delivers the real message to a few named addresses, through the real gate, without touching the segment. Do this every time.
- 8 **Submit.** Below the approval threshold it approves itself. Above it, it goes to an Approver.
- 9 **Send now, or schedule.** Both need `MANAGER`. A confirmation dialog states exactly who is about to be messaged and which exclusions are in force.

Reading the audience line

Under the segment picker you will see something like:

24,929 will receive • 25,100 in segment • 33 suppressed/opted-out • 138 excluded by your rules

in segment	Everyone the segment rule matches, before any filtering.
suppressed/opted-out	Dropped because they have a hard block, or have not opted in to marketing on this channel.
excluded by your rules	Dropped by the exclusions you set on this campaign.
will receive	The number that matters. This is what the approval threshold is judged on, and what the confirmation dialog repeats back to you.

Send test

A test send uses the campaign's real template and goes through the real send gate, so what lands on your phone is what a customer would get. It is recorded separately and never appears in the campaign's own numbers. You can run it as many times as you like.

 IMPORTANT

When test mode is off, **Send now** reaches real customers immediately. The confirmation dialog will say so in as many words. Read it rather than clicking through it: it names the audience and lists the exclusions, and it is the last point at which stopping is free.

What happens after you press Send

Sending is no longer a single leap. The moment you press **Send now** (or the schedule fires), the campaign first takes a snapshot of exactly who it intends to message — the *roster* — and only then begins delivering, in several parallel lanes. Two consequences worth knowing:

Preparing the audience	For a short while the campaign shows this instead of sending. It is normal — a big audience takes a minute or two to snapshot — and sending starts by itself when it finishes. Stop works during this phase too.
N never attempted	After the campaign finishes, this badge is an <i>exact</i> count of people who were on the roster and never got any attempt — not sent, not failed, not skipped. One press of Send to the missed reaches exactly them; everyone already handled is skipped automatically.
Add people who became eligible since	New opt-ins keep arriving after the snapshot. This button tells you how many, and adds them to the roster — nothing sends until you press Resume.
Stalled	If a delivery lane fails repeatedly, the campaign parks itself here and posts an alert, instead of silently looking busy forever. Resume continues exactly where it stopped.

i WHY THE NUMBERS CAN DIFFER FROM THE AUDIENCE LINE

The audience line is live and keeps moving as people opt in and out. The roster is frozen at the moment you sent. “Who did we miss?” is always answered against the roster, so the answer stays the same no matter when you ask.

Audience Exclusions

/campaigns

AUTHOR

MANAGER

APPROVER

Exclusions let you say “send to this group, but not to these people”. They are the answer to the most common broadcast problem: the same customer getting two things from us in one day.

WHERE	THREE RULES	WHEN CHECKED	WHO
On the campaign form, under the segment picker.	All optional. Anyone matching <i>any</i> of them is skipped.	Continuously, while the campaign sends.	AUTHOR MANAGER

The three rules

Don't send to these segments	Tick any number of other segments. Anyone in them is skipped. Use it to hold back staff, existing buyers, or a group you messaged separately.
Already reached by these campaigns	Tick previous campaigns. Anyone who received one is skipped. This counts across every channel, because a campaign only has one channel: excluding a WhatsApp blast will correctly hold those people back from a follow-up email.
Recently contacted on this channel	Skip anyone we have messaged on this channel within a chosen window, from 6 hours up to 30 days. Unlike the rule above, this one is same channel only : an email this morning will not stop a WhatsApp tonight.

Why the timing matters

A large broadcast takes hours to work through its audience. Exclusions are re-checked continuously as it goes, not fixed at the moment you pressed send. So if a colleague sends something else while your campaign is still running, the people who received theirs are dropped from the rest of yours automatically.

Two campaigns running at the same time will also hold each other back, so you can start both without co-ordinating them by hand.

What counts as “contacted”

Only messages that genuinely went out. A message the send gate stopped, for example because the person had hit the frequency cap, does **not** count: they received nothing, so excluding them would be

wrong. Journey messages do count, because from the customer's side a journey message is just as real as a campaign one.

✓ **TIP**

Set the rule, then read the audience line. It updates as you tick, and tells you exactly how many people each rule removed before you commit to anything.

① **NOTE**

Exclusions lock when you submit the campaign, along with everything else. The confirmation dialog on **Send now** lists every rule in force, so a rule you set days ago and forgot cannot silently shrink your send.

A/B Testing

/campaigns

AUTHOR

MANAGER

Split a campaign into arms, send each to a slice of the audience, and let the numbers say which message worked. Set up on the campaign itself, read afterwards in the Experiment log.

WHAT IT'S FOR	WHERE	WHO	LOCKS AT
Proving which message performs better.	The Variants panel on a draft campaign.	AUTHOR MANAGER	Approval. Arms cannot change once approved.

Setting up a test

- 1 **Open a draft campaign.** Variants can only be added while it is a draft.
- 2 **Add an arm.** Each arm gets a label (A, B and so on), a template, and a weight that decides what share of the audience it gets.
- 3 **Check every arm is sendable.** On WhatsApp, every arm's template must be approved by Meta. The campaign refuses to start otherwise, rather than sending half a test.
- 4 **Send as normal.** Assignment happens per person as the broadcast runs.

Holdout arms

An arm with no template attached is a **holdout**: those people are deliberately sent nothing. It is how you measure whether the message did anything at all, rather than which of two messages did more. If the holdout converts as well as the treatment, the campaign was not the reason.

⚠ WATCH OUT

A holdout arm is people who receive **nothing**, on purpose. Do not add one to a campaign carrying information customers need, like a delivery change or a price correction.

How people are assigned

Each recipient is assigned to an arm as they are reached, not in blocks. This matters more than it sounds: if all of arm A went out first and all of arm B hours later, the test would be measuring time of day rather than the message. Assignment is also stable, so the same person always lands in the same arm for a given campaign.

① **NOTE**

Results live in the Experiment log, not here. This screen is for setting the test up; that one is for reading it and writing down what you learned.

Experiment Log

/experiments

VIEWER

AUTHOR

MANAGER

Every A/B test Relay has run, in one list, with the result of each. It is a log across all campaigns rather than a screen you build anything on.

WHAT IT'S FOR	READ OR ACT?	WHO	COVERS
Remembering what you already learned.	Read-only, except for recording a learning.	Anyone with Relay access can read it.	Every campaign that had more than one arm.

What each row shows

Campaign	Which broadcast the test ran on. Click through to open it.
Arms	Each arm, how many people it reached, and how it performed.
Learning	A short note you write yourself once the result is in. This is the part that makes the log worth having.

Reading a result honestly

Two numbers being different is not the same as one message being better. A few hundred people per arm can produce a gap that is entirely noise. Treat a small difference on a small audience as “no result”, and be willing to run it again on something bigger.

✓ TIP

Write the learning in plain words and include what you would do next time: “shorter subject line won on a 4,000 send, worth repeating on the next launch” is useful in six months. “B won” is not.

Journeys

/journeys

VIEWER

AUTHOR

MANAGER

A journey is an automation. Something happens (a cart is abandoned, an order ships), and Relay reacts, on its own, for that one customer. Journeys run continuously, which makes them the highest-leverage and highest-risk thing in Relay.

WHAT IT'S FOR	READ OR ACT?	WHO	RUNS
Reacting to customer behaviour without anyone pressing send.	Action screen. An active journey sends by itself.	AUTHOR builds, MANAGER activates.	Continuously, one enrolment per customer.

Working with nodes on the canvas

Add	The buttons above the canvas add a step of that type. Drag from a dot on a node's right edge to another node to connect them.
Replicate	Click a node, then Replicate in the panel below. You get a copy with the same settings, offset slightly. The copy is not wired to anything — connect it yourself.
Delete	Delete node in the same panel, or select the node and press Backspace or Delete. Deleting also removes the connections into and out of it.

Replicate is for building a branch that differs in one detail — a second send with different copy, a second wait with a different delay. Copying rather than rebuilding keeps the settings you already got right.

NOTE

Both actions only appear on a journey you can edit. Open an *active* journey and the canvas is read-only, so neither button shows — pause it first, or work on a draft.

The parts of a journey

Trigger	The event that starts it, for example <code>checkout_abandoned</code> . Optionally narrowed by a filter, so it only fires for certain values.

Steps	The graph the customer walks through: send a message, wait, branch on what they did, send again.
Enrolment	One customer's trip through the journey. Each has its own position and its own timers.
Version	A frozen copy of the steps. An enrolment stays on the version it started on, so editing a journey never rewrites the path someone is already walking.

Saving: what gets checked

Relay checks the whole graph every time you save, and it refuses to save a journey that could not run. This is deliberate: a journey that saves with a hole in it can be switched on later by someone who did not build it, and the hole only shows up as customers silently not receiving messages. The checks you are most likely to meet:

Send step with nothing to send	Every send step needs <i>either</i> a saved template <i>or</i> message text typed on the step itself. With neither, the step has no message and every send from it would fail. Shown as <code>send_no_template_or_body</code> against the step name.
A branch that goes nowhere	Every outcome on a step has to point at another step. A condition needs both its yes and no branches wired, and a reply-button step needs each button wired.
Template not active	A send step pointing at a draft or archived template will not save. Get the template approved and active first.
A wait you cannot read	Wait durations are read as plain text, so <code>3 days</code> works and a typo like <code>3 dayz</code> is refused rather than quietly treated as zero.

Failures appear above the canvas as **Validation errors**, naming the step so you can go straight to it. Fix the step and save again.

⚠ WATCH OUT

Free-text steps are allowed and are not a mistake. A step with typed message text and no template is exactly how the reply parts of the COD confirmation flow work. The check is only that a step has one of the two, never that it must use a template.

Trigger filters

A filter narrows the trigger to certain property values, with `=` or `≠`. The `≠` option exists for a reason worth understanding: excluding a category by saying "equals everything else" silently drops every event

where the property is missing altogether, and on some events that is a large share. If you mean “not this one”, say .

Re-enrolment

Each journey decides what happens when a customer triggers it again while already inside it. The common setting is one active enrolment at a time, so a customer who abandons two carts in an hour is not messaged twice. Some journeys instead re-enrol every time, when each trigger genuinely deserves its own message.

WATCH OUT

Deactivating a journey stops new enrolments *and* stops customers already inside it. It is a real stop, not a “let the current ones finish”. That is usually what you want in an emergency, but be aware that half-finished journeys will not resume if you switch it back on.

Exclusions

A journey carries the same three exclusion rules a campaign does, set once in Settings for the whole journey: don’t send to chosen segments, don’t send to anyone already reached by chosen campaigns, and don’t send to anyone contacted recently on the same channel. They are checked *each time a message is about to go out*, never when someone enrolls, so a customer who receives a broadcast on day two of a five-day journey is held back from the steps that follow.

A person the rules match is skipped for that message only — they stay in the journey and carry straight on to the next step, exactly as they would after any other skip. Nobody is ever removed from a journey by an exclusion.

Quiet hours inside a journey

When a journey step lands in quiet hours, the message is held and retried at the boundary rather than dropped. A customer whose order ships at 23:00 gets the message in the morning, not never.

IMPORTANT

Never edit a template that a live journey depends on. The journey is pinned to the template, and changing the wording changes what is already going out, with no review step and no announcement. Duplicate the template, edit the copy, and point the journey at it deliberately.

Bot Builder

/journeys

VIEWER

AUTHOR

MANAGER

A bot is a scripted conversation on the website chat widget (and optionally WhatsApp): menus, order status, agent handoff. It uses the same canvas as journeys, on a **Journeys | Bots** toggle. Bots have no trigger, no exit rules, no versions screen and no funnel — just a definition, a status, and a test.

WHERE	READ OR ACT?	WHO	CHANNELS
Journeys screen, Bots toggle.	Action screen. A published, active bot answers customers on its own.	AUTHOR builds and saves drafts, MANAGER also publishes, pauses/resumes and sets WhatsApp rollout.	Web widget, WhatsApp (support number), or Shared flow (used by other bots).

The bot list

Each row is a bot: name, status (draft / active / paused), channel, live version, and its last 7 days — chats, handled, handoffs to agents, and conversions (a session whose customer placed an order within 24 hours of the chat starting). **Open** loads it onto the canvas; **New bot** starts a blank web-widget draft.

Building the flow

Add a step	Buttons above the canvas add: Message, Menu, Ask for input, Order status, Hand to agent, End chat, Shared flow . Drag from a node's right-edge dot to another node to wire an outcome.
Chat start	The fixed anchor node (the same node journeys use for their trigger, relabelled). It fires when the customer opens the widget (on WhatsApp: their first message); wire its one outcome to your first step.
Duplicate step	Select a node, then Duplicate step in the panel below. You get a copy with the same settings, offset slightly. The copy is not wired to anything — connect it yourself. The copy becomes the selected node, so the panel opens on the new one, not the original. A copied Menu keeps its option labels but gets new option ids, so a customer tapping an old button on the original menu is never sent down the copy's branch.

Delete step	Delete step in the same panel, or select the node and press Backspace or Delete. Removes the connections into and out of it too.
Expand	The Expand button (top-right of the canvas) fills the screen, docking the Step / Settings / Test panel beside it. Esc exits.

Only one connection is kept per outcome: dragging a new wire from a dot that is already wired replaces the old one, it does not add a second.

Step types

Message	Sends text, one outcome (<code>next</code>).
Menu	Shown as Buttons or a List . One outcome per option, plus Not understood (the <code>fallback</code> branch, taken after two misses in a row — wire it, usually to Hand to agent). Matches either a tapped button or a typed reply that matches an option's label or number; anything else re-shows the menu until the miss cap.
Ask for input	Collects Phone or email (identity) or an Order number . An invalid answer re-asks once; a second invalid answer takes Not understood (the <code>fallback</code> branch — publish refuses it unwired). Typing a keyword configured in Settings jumps straight to that keyword's step instead of counting as a miss.
Order status	Needs a Phone or email ask and an Order number ask earlier in the flow. Looks the order up and verifies it: the order's own phone/email must match what the customer already gave. Order found and Not found / no match are the two outcomes — a mismatch and a genuinely missing order both take Not found / no match , on purpose, so the bot never confirms which order numbers exist. After 5 failed lookups the bot says the step's configured text and hands off.
Hand to agent	Handoff text is optional and editable — type what the bot should say as it hands over. Leave it empty and it uses the standard wording ("Let me connect you to our support team — a human will reply right here as soon as one is available."). Terminal. Ends the bot's part of the conversation and lands it in the Pitstop inbox as a thread; the bot never speaks again in that conversation. The order-lookup cap hands off the same way — a customer the bot cannot help reaches a person, never silence.
End chat	Terminal. Optional goodbye text, then the session ends.

Shared flow

Jumps into another bot built with channel **Shared flow**. Only a shared flow that is *published and active* can be picked. When the shared flow ends, the customer continues on this step's `next`. The chat runs the shared flow's live version *at the moment it enters* — so publishing a shared flow changes it for chats already under way, and pausing one hands those chats to an agent when they reach this step.

⚠ WATCH OUT — WHATSAPP LIMITS

Buttons: max 3 options, 20 characters each. List: max 10 rows, 24-character titles. The canvas only warns about these (it does not know the bot's channel yet); **Publish enforces them for real**, per channel — including inside a Shared flow, since it can be called from a WhatsApp bot too.

Settings

Channel (fixed after the first publish, since sessions and the inbox key on it), keywords (an answer the bot did not accept jumps straight to a step; on WhatsApp the opening message counts too — the web widget opens without one), and for a WhatsApp bot, rollout: **Pilot** (only the numbers listed) or **Public** (every customer on the number). Rollout needs the activate permission, and going Public is Afshaan's call.

Test panel

Runs the current draft (unsaved edits included) against a simulated customer. It is side-effect-free: nothing is written, no order is looked up and nobody is handed off — those steps show as → *would look up...* and → *would hand off to an agent here*. Entering a Shared-flow step in the test always follows that flow's **draft**, not its published version — the live bot always runs the published one, so a draft-only difference in a shared flow will not show up here.

NOTE

The Test panel only appears when nothing else is selected. Click empty canvas to get back to it after configuring a step.

Saving and publishing

Every outcome on every step must be wired before you can publish — an unwired menu button or an unwired fallback is a customer reaching a dead end. **Save draft** is available to anyone who can build; **Publish** (freezes a version; every new conversation gets it at once) needs the activate permission.

NOTE

Publish saves the canvas first, then checks and publishes it — what you see (and just tested) is what goes live. If the save fails, nothing is published.

⚠ TWO PEOPLE ON ONE BOT

If someone else saved this bot (or paused, resumed or published it) after you opened it, **Save draft** and **Publish** stop and tell you so instead of silently replacing their work. You can **Cancel** — your edits stay on screen, nothing is saved, and you reload the bot to see theirs — or type **OVERWRITE** to replace their version with yours. An unpublished draft has no history, so an overwrite cannot be undone. If a colleague saves in the moment between your save and your publish, nothing is published: reload and review first. A tab left open from before this check existed is asked to reload.

CODE	WHAT IT MEANS	FIX
no_entry	Chat start is not connected to a first step.	Wire Chat start to your first step.
button_unwired	A menu option has no outcome wired.	Connect that option to a step.
fallback_unwired	A menu or Ask for input step has no Not understood / Fallback branch wired.	Wire it, usually to Hand to agent.
dangling_target	A step's outcome is not wired (a fresh Duplicate step copy, for one), or points at a step that was deleted.	Wire the named step's outcome to a step.
menu_no_buttons	A Menu step has no options at all.	Add at least one option.
subflow_target_invalid	A Shared flow step points at a bot that is not a published, active shared flow.	Pick a currently-published shared flow, or publish the one you meant.
shared_contains_subflow	A Shared flow bot itself contains a Shared flow step.	Shared flows cannot nest — build those steps in directly.
keyword_empty / keyword_target_missing	A Settings keyword has no words, or points at a step that no longer exists.	Fill in the words or re-pick the step in Settings.

CODE	WHAT IT MEANS	FIX
wa_too_many_buttons / wa_button_label_long	A WhatsApp-facing menu (this bot's channel, or a shared flow called from one) exceeds 3 buttons or 20 characters.	Trim the options or switch to a List.
wa_too_many_rows / wa_row_title_long / wa_row_description_long	A WhatsApp-facing list exceeds 10 rows, or a title/description is too long.	Trim rows or shorten the text.
wa_text_long / text_exhausted_long	A message body (or the Order status exhausted-attempts text) is over WhatsApp's length limit.	Shorten the text.

Pause and resume

Pause keeps the published version but stops new conversations. On the **web widget**, a new visitor sees *Chat is unavailable right now — please email support@legendoftoys.com*; a chat already open carries on to its end. On **WhatsApp** the bot goes quiet at once, even mid-chat — a customer it was talking to can sit out of the Pitstop *Awaiting* list while they keep writing, so check the inbox for them after a pause. Pausing a **Shared flow** hands off any chat that reaches it. Resume — or Publish, which also sets the bot active — brings it back.

IMPORTANT

Save draft never touches what customers see. **Publish** does, at once, for every conversation that starts after it — a conversation already in progress finishes on the version it started with, except inside a Shared flow, which always runs that flow's current version. Run the change through the Test panel before you publish.

Link Tracking & Attribution

/campaigns VIEWER AUTHOR MANAGER

Whether a campaign can be measured is decided *before* it is sent — in fact before the template is even submitted to Meta. This chapter explains why, and what to do so a send is never wasted.

WHAT IT'S FOR	DECIDED WHEN	WHO	WHERE
Making clicks and revenue countable.	Before Meta approves the template.	AUTHOR MANAGER	Templates → Buttons, and Links.

 THE ONE THING TO REMEMBER

A WhatsApp button's web address is **frozen when Meta approves the template**. Nothing can be added to it when the message is sent. So filling in UTM values on a campaign does *nothing* if the button is a plain link — and there is no way to fix it afterwards. It has to be arranged first.

What happened once, so it does not happen again

An Independence Day campaign went out to several thousand people with UTM values filled in on both the campaign and the template. It recorded **zero clicks and zero attributed revenue**, and always would have. The button was a plain link, so nothing in the system could tag it. Nobody had done anything wrong — the screen simply never said so. It says so now, and it will stop you.

The three states a button can be in

STATE	WHAT YOU GET	WHAT TO DO
Tracked per person	Who clicked, and what they bought	Nothing — this is the goal
Tracked, shared link	Total clicks for the campaign	Fine. You will not know <i>who</i> clicked
Not tracked	Nothing at all	Fix before sending — see below

Option 1 — track every person separately (best)

Do this once per template and every future campaign using it is tracked.

- 1 When you submit the template to Meta, set the button address to `https://lottoys.in/r/{{1}}` instead of the real destination.

- 2 Once Meta approves it, open the template in **Templates**, find the button, and tick **Track clicks on this button**.
- 3 Put the real destination — where people should actually land — in the field that appears.

Everyone then receives their own link, so clicks and orders can be traced back to individual people.

⚠ DO NOT SKIP STEP 2

A template approved as `/r/{1}` but never ticked is *worse* than an untracked one — the button can send people to a dead link. Relay will warn you if it spots this, but do not rely on the warning.

Option 2 — one shared tracked link (quicker)

Useful when you want the campaign out and only need total clicks.

- 1 In the campaign, when Relay warns the links cannot be tracked, choose **Create tracked link**. You can also make one directly under **Links**.
- 2 Give it a short name, the destination, and any UTM values.
- 3 Copy the address it gives you and use *that* as the button address in your Meta submission.

Once the approved template carries that address, its clicks appear on the campaign automatically. There is nothing else to connect up. The link also appears under **Links**, where the destination can be changed later.

What Relay does to stop you

Picking a template whose links cannot be tracked opens a warning straight away — not at the end, when the copy is written and the send is imminent. The same warning is repeated in the confirmation box when you press send. Neither one blocks the campaign: you can send an untracked campaign deliberately. You just cannot do it by accident.

✓ TIP

Whatever the link situation, a discount code unique to the campaign is worth adding. It survives every tracking gap and lets Finance attribute revenue even when the click data is missing.

Email is different

Email bodies are written here rather than approved elsewhere, so Relay tags links in an email automatically. UTM values on an email campaign work exactly as they appear. This whole chapter is about WhatsApp, where the button belongs to Meta.

PART 4

Audience

Who Relay knows about, what they have done, and how to carve them into groups.

Activity

/activity

VIEWER

AUTHOR

MANAGER

ADMIN

A live feed of what customers are doing: pages viewed, carts abandoned, orders placed. It is the raw material every segment and journey is built on.

WHAT IT'S FOR Seeing the events arriving, and confirming they arrive at all.	READ OR ACT? Read-only.	WHO Anyone with Relay access.	RANGE Pick a window. Defaults to today.
---	----------------------------	----------------------------------	--

The three counters

Identified	Events we could attach to a known person. An event from an anonymous browser counts only once we learn who they are.
Messaged	Of those, how many we then sent something to.
Recovered	Of those, how many went on to buy.

Why this screen is worth checking

Events are the input to everything else. If a feed quietly stops, segments keep returning their last answer and journeys simply stop firing, with no error anywhere. A glance at Activity is the cheapest way to catch that: if a familiar event has gone to zero today, something upstream has broken.

NOTE

Not every event reaches back to the beginning of time. Website and order events only exist from the day their feed was switched on. Counting a customer's lifetime orders from events will undercount badly; use the customer's stored order count instead. The Segments chapter covers this.

Segments

/segments

VIEWER

AUTHOR

MANAGER

A segment is a rule that picks out a group of customers. Everything you send is aimed at one, so this is where most of the thinking happens.

WHAT IT'S FOR Defining who a message is for.	READ OR ACT? Action screen, but it sends nothing.	WHO AUTHOR MANAGER	RECOUNTED On Refresh members, and on every send.
---	--	-----------------------	---

How a rule is built

A segment is a group of conditions plus a matching mode:

MODE	MEANS
Match ALL of	Every condition must be true.
Match ANY of	At least one must be true.
Match NONE of	Everyone for whom none are true.

Three kinds of condition

Attribute	Something we know about the person: <code>lifetime_orders</code> , <code>lifetime_value</code> , <code>last_order_at</code> , <code>city</code> .
Event	Something they did, counted, optionally inside a time window and optionally narrowed to a specific value.
Consent	Their opt-in state on a channel.

The attribute operators

OPERATOR	MEANS
is exactly / is not exactly	Equal to, or not equal to, a single value.
is any of	Matches any value in a list.

OPERATOR	MEANS
more than / at least	Greater than, or greater than or equal to.
less than / at most	Less than, or less than or equal to.
older than (days)	For a date: longer ago than N days.
within last (days)	For a date: in the last N days.

"Bought exactly once" is `lifetime_orders` **is exactly** `1`. "Never bought" is `lifetime_orders` **is exactly** `0`.

Counting events

An event condition reads "`order_placed` $\geq$ 1 within last 90 days". The operator can be $\geq$, = or $\leq$, so "did this exactly twice" and "has never done this" are both sayable. The **where** box narrows to a specific value, for example only views of one collection.

⚠ WATCH OUT

Counting orders from the `order_placed` event is not the same as the customer's real order history. Events only exist from the day that feed was switched on, so an event count reaches back weeks, not years, and will badly undercount long-standing customers. For anything about lifetime buying behaviour use the `lifetime_orders` attribute instead.

Grouping conditions

Add group puts a bracket inside your rule. Everything in the group is worked out first, then the result joins the rest of the list under the mode at the top. That is how you say something like *bought at least once, and (lives in Pune or lives in Mumbai)*: set the top to Match ALL, add the purchase condition, then add a group set to Match ANY holding the two cities.

A group can hold conditions but not another group. One level covers what people actually need, and deeper nesting gets hard to read at a glance.

✓ TIP

A new group starts as Match ANY, because that is almost always why you are opening a bracket. A group set to the same mode as the level above it does nothing.

① **NOTE**

An empty group is ignored when you save, so an unfinished one cannot quietly change your audience. If you see a rule shown read-only with a note about groups inside groups, it was written with more nesting than this screen can display: it still works, and it is locked so that editing cannot flatten it.

Members, and why the count can read zero

A segment's member list is rebuilt when you press **Refresh members**, and again automatically whenever it is used to send. A segment nobody has ever refreshed shows **Not counted** rather than a number, because zero and unknown are very different things.

🚫 **IMPORTANT**

An incomplete condition row will be refused when you save, rather than quietly matching nobody. If you see an error naming a row, finish that row: an empty condition used to drag an entire segment to zero members, which is indistinguishable from "there are no such customers" and has caused a real send to go nowhere.

✓ **TIP**

Use the preview and the sample rows before trusting a new rule. If the sample shows people you did not expect, the rule is wrong, and it is much cheaper to find out here than in Analytics afterwards.

Contacts

/contacts

VIEWER

AUTHOR

MANAGER

ADMIN

One customer at a time: everything Relay knows about them, everything it has sent them, and why they might not be receiving anything.

WHAT IT'S FOR	READ OR ACT?	WHO	SEARCH BY
Answering "why did this person get this?" or "why did they get nothing?"	Mostly read. Consent edits need ADMIN .	Anyone with Relay access can read.	Name, email or phone, including partial.

The contact detail panels

Profile	Name, city, and the stored facts a segment can filter on, such as lifetime orders and last order date.
Identifiers	Every email and phone number we have linked to this person. One profile can carry several.
Consent	Their opt-in state per channel, with the history of how it changed and where each change came from.
Delivery blocks	Hard blocks. This is the panel that explains total silence: a permanent bounce or an SMS opt-out lands here and stops everything.
Recent events	What they have done lately.

Consent and opt-out

Consent is a running log, not a switch. Every change is appended with its source, so you can always see when someone opted in and where that came from. **Opt out everywhere** records a withdrawal across all channels.

⚠ WATCH OUT

An opt-out and a delivery block are different things. Opting someone out of marketing still lets their order updates through, which is correct. A delivery block stops *everything*, including order updates, so only lift one if you know the address is genuinely working again.

Anonymous profiles

Many profiles are anonymous browsers we have not yet matched to an email or phone. They are real records and they do count in totals, but they cannot be messaged. If a segment's total looks much larger than the number that will actually receive a campaign, unmatched anonymous profiles are usually the reason.

Forms

/forms

VIEWER

AUTHOR

MANAGER

ADMIN

The storefront capture forms Relay knows about — who is signing up, and for the back-in-stock form, whether they have been alerted yet.

WHAT IT'S FOR	READ OR ACT?	WHO	SEARCH BY
Seeing who is signing up on a <code>/f/*</code> form, without SQL.	Read only. Nothing on this screen sends anything or changes a journey.	Every Relay viewer — contact details are visible here at <code>VIEWER</code> , same as Contacts.	Name, email, phone or SKU, on the sign-ups page — filters the current page only.

What a form is

A form is one of the storefront capture widgets at `/f/*`. The **back-in-stock** form is embedded on sold-out product pages via `data-lot-form="back-in-stock"`; a customer who wants to know when a sold-out variant returns submits it there. A form appears on this screen as soon as it exists in `comms.forms` — there is nothing to configure to make it show up.

The Forms list

One row per form: its name and slug, **Kind**, then **Total**, **7 days** and **30 days** sign-up counts, and the timestamp of the last sign-up. **Open** goes to that form's sign-ups.

NOTE — WHAT ACTIVE MEANS

Active means the form is accepting sign-ups. It is **not** a sending switch: whether anyone is ever messaged is decided separately, by the journey that listens for the form's event. A form can be active for months with its journey still in draft, collecting sign-ups that nobody has been told about yet.

A dash (—) in any count column means the count could not be read, not that it is zero.

One form's sign-ups

Opening a form shows its tiles, a 30-day chart, and the sign-up table.

Tiles	Total sign-ups, distinct contacts, last 7 days, WhatsApp opt-ins, and for the back-in-stock form, Alerted.
-------	--

Sign-ups per day	A 30-day bar chart by IST calendar day. Hover a bar for that day's count.
Top products	Back-in-stock only. Ranked by sign-ups per Shopify variant SKU, with the LOT product code shown where <code>sales.sku_map</code> knows it. A handful of early test sign-ups stored no SKU and are counted separately as "No SKU".
Table	Signed up (IST), Name, Email, Phone, Product (back-in-stock only), Channels, and Alert (back-in-stock only). Name, email and phone are what the customer typed on the form itself, falling back to their contact profile's newest identifiers when the form left one blank.
Search & paging	Search matches name, email, phone or SKU on the current page only — it does not requery the server. 100 rows per page; Newer / Older move a page at a time. The tiles and charts always cover the whole form, not just the visible page.

▲ BEYOND 5,000 SIGN-UPS

Total, 7-day and alerted counts are always exact. Distinct contacts, the per-day chart, opt-in counts and top products are computed over the newest 5,000 sign-ups; past that size the page adds a note that those figures are sampled, not the whole population.

Back-in-stock alerting

🚫 IMPORTANT

Sign-ups are recorded even while the back-in-stock journey is **DRAFT**. No customer is emailed until the journey **"Back in stock — requested alert"** is activated. Each sign-up can be alerted at most once, ever — the ledger enforces one alert per submission, so a product going out of stock and back three times still sends one email, not three. **While the journey is draft, no sign-up is used up.** When it is switched on, a waiting sign-up is alerted only if its product is *still* in stock at that moment; if it came back and sold out again while the journey was draft, it simply waits for the next restock — there is no burst of stale "back in stock" emails on go-live day.

The **Alert** column reads the same ledger the sending job writes to:

STATUS	WHAT IT MEANS
Waiting	Not back in stock since this sign-up, or the journey is not active yet. The normal state for most rows.
Queued	Back in stock — the alert is about to fire.
Retrying	The alert failed at least once and will be tried again.

STATUS	WHAT IT MEANS
Stuck	The alert failed five times in a row and is no longer retried automatically. Hover the badge for the last error.
Alerted	The back-in-stock event fired for this sign-up.

Who can see it, and what is missing

Every Relay viewer can open Forms and read a sign-up’s email and phone — this screen is gated the same way Contacts is, not more tightly, because it is the same class of customer data. There is no CSV export yet; if you need the list outside Relay, that is a request to raise, not a workaround to build.

PART 5

Build & Measure

The things a send draws from, and the numbers that tell you how it went.

Templates

/templates

VIEWER

AUTHOR

MANAGER

A template is the message itself. Campaigns and journeys both point at one rather than carrying their own copy, so a template is shared and editing it changes everything using it.

WHAT IT'S FOR Writing the message, once, per channel.	READ OR ACT? Action screen.	WHO AUTHOR MANAGER	VERSIONED Yes. Every save keeps the previous version.
--	--------------------------------	-----------------------	--

Channels behave differently

Email	A subject line and a body, built in the visual canvas. No outside approval needed: write it and use it.
WhatsApp	Must be submitted to Meta and approved before it can be used for a broadcast. Approval takes time and can be refused. Plan for it.
SMS	The wording must be registered with the Indian DLT registry first (the one step outside Relay); then the editor's Create at TrustSignal button pushes it to the sending vendor in one click. What goes out has to match the registered wording exactly.
RCS	Composed and submitted entirely in Relay — write it, press Submit to TrustSignal , and Sync pulls the approval. Every RCS template also names an SMS fallback template, because phones that can't receive RCS get the SMS instead. Full walkthrough in the SMS & RCS chapter.

Variables

A variable is a placeholder filled in per recipient, like a first name or an order number. Every variable must resolve for that person, or the send is skipped rather than going out with a gap in it. Use the preview to confirm each one fills.

Duplicating instead of editing

Duplicate makes an independent copy. Reach for it constantly. Editing a template that a live journey or an approved campaign depends on changes what is already going out, with no review and

no announcement.

 IMPORTANT

Never edit a template a live journey depends on. This has reached real customers before. Duplicate it, edit the copy, then repoint the journey deliberately.

Archiving

Archiving retires a template from the pickers so nobody wires it up again. Anything already using it keeps working, on purpose: silently breaking a live flow would be worse than letting an old template continue.

 WATCH OUT

Deleting a WhatsApp template on Meta's side is not reversible from here and removes every language version at once. Archive in Relay instead unless you are certain.

Warnings when you save an email

The email builder checks your template every time you save. If anything is wrong with the underlying markup, a yellow panel appears under the editor listing what it found.

These do not stop the save, and they are worth reading anyway. The checker only warns — it will happily save a template it knows is broken. A real example: a fault in one template's global styling meant none of its font settings applied, so it went out to customers at the wrong size. The checker had spotted it every single time; nobody could see what it was saying. Now you can.

If you see warnings you do not understand, fix them before setting the template Active rather than after — once it is live the next thing that reads it is a customer's inbox.

Choosing a purpose

The **Purpose** dropdown decides which rules a send has to pass — whether it needs marketing consent, whether it obeys quiet hours, and which number it goes out from. Pick the one that describes the message honestly; it is not a label, it is the rule set.

SMS & RCS

/templates

VIEWER

AUTHOR

MANAGER

SMS and RCS are the phone-message channels. SMS is the plain text message every phone can receive. RCS is its richer successor — an image, a card and a tappable button, delivered through Google Messages — and when a phone can't receive RCS, the message automatically goes out as a plain SMS instead. You author both in Relay; nothing needs a second portal except one registration step for new SMS wording.

WHAT IT'S FOR	READ OR ACT?	WHO	AUDIENCE
Reaching customers off WhatsApp and email, with automatic SMS fallback.	Action — authored on the Templates screen, sent from Campaigns and Journeys.	AUTHOR MANAGER	The SMS opt-in list — roughly a tenth the size of WhatsApp's. A reliability channel, not a reach one.

The rule that shapes everything: sending hours

Indian carriers only deliver promotional SMS and RCS between **10:00 and 21:00**. This is the carrier's rule, not ours — a message pushed outside the window is refused by the network *and still costs credits*. Relay knows this: a campaign scheduled outside the window parks and goes out when the window opens. Don't fight it; plan sends for mid-morning.

Writing an RCS template — entirely in Relay

- 1 New template → channel RCS → Save.** Saving first matters: the approval is stamped onto the saved row.
- 2 Compose.** Write the message in the Compose & submit box. Put variables in `[square_brackets]` — `[name]`, `[code]`, `[link]`. For a rich card, add the card title and an image URL (upload the image to the Library first and copy its address; landscape 2:1, under 2 MB).
- 3 Button (optional but worth it).** A button needs all three fields — text, URL and an id. Missing the id is the one thing the carrier rejects outright.
- 4 Submit to TrustSignal.** One click. Plain text templates usually approve in about a minute; rich cards take longer. Press **Sync status** to pull the verdict — the badge turns green on approval, and a rejection shows its reason in red.
- 5 Pick the SMS fallback.** Every RCS template must name an approved promotional SMS template — that is what a non-RCS phone receives. The picker only offers valid ones.
- 6 Declare the variables.** One entry per bracketed name, exactly as on any other channel.

Writing an SMS template — one outside step, then Relay

- 1 **Register the wording on the DLT portal.** This is the legal registration Indian SMS requires and the only step outside Relay. It produces a 19-digit id. Approval is on the operator's clock — never same-day-of-send.
- 2 **Author in Relay.** New template → SMS. Enter the registered wording with each placeholder as a named `{token}`, the variable order, the consent type, and paste the 19-digit id.
- 3 **Create at TrustSignal.** The button arms only once the id is in place — one click pushes everything to the vendor. No retyping into any portal.

Already made it in Sigo? Adopt it instead

If the template was written at the vendor first, do *not* retype it here — press Find templates in Sigo on the SMS panel and press Adopt next to it. Relay copies the registered wording, the consent type, the DLT id and the placeholder count straight off the vendor, so what you send is exactly what was registered.

- Nothing listed? Every Sigo template is either already adopted here, or carries a consent type this system cannot send on. Only `Service-Explicit` and `Service-Implicit` can send — a template made as `Promotional` is refused at adopt rather than failing later in front of a customer.
- After adopting, name the placeholders in Variable order. Until you do, the template holds the raw `{#var#}` markers and will refuse to send — deliberately, so a marker can never reach a customer as literal text.

IMPORTANT

Never reword a registered SMS body — not a comma. The carrier compares the delivered text against the registration and silently refuses a mismatch. Only the `{token}` values change. New wording means a new DLT registration.

What the numbers will show

Delivered	Both channels report it, per message.
Read	RCS only. SMS has no read receipt — an opens column showing 0 for SMS is honest, not broken.
Clicks	Both channels. With the template's tracked-link fields set, every recipient gets their own short link, so you see <i>who</i> clicked, not just how many.
Fallback	An RCS message that went out as SMS is counted on the SMS side of the analytics — one message, one row, whichever leg delivered.

Test sends

The test box takes a phone number (10 digits; SMS is India-only so +91 is fixed). Tests reach approved test numbers only, work at any hour, and skip the marketing gates — but the *carrier* still enforces its own sending hours on promotional routes, so a 2 a.m. test can show as failed with a “time” error. That is the carrier, not the template.

⚠ WATCH OUT

A campaign to an RCS audience reaches the people opted in to SMS marketing — check the **will receive** number before promising reach. It is a fraction of a WhatsApp audience, and that is expected.

Correcting an RCS vendor binding

An RCS template carries the id of the creative registered at the vendor. Once set, the field is read-only — a stray keystroke there would silently point the template at something that does not exist.

If the id is *wrong*, click **Unbind** beside it. The field clears and becomes editable, so you can either type the corrected id and save (a straight rebind), or save it empty to leave the template unbound.

⚠ WATCH OUT

Unbinding only takes effect when you **Save**. It is refused if the template is attached to a live campaign, and an *active* RCS template cannot be saved without a binding at all — set its status back to draft first.

Library

/library

VIEWER

AUTHOR

MANAGER

Every image available to an email or WhatsApp template, in one place. Upload here once, then pick from the library while writing.

WHAT IT'S FOR Storing and reusing images.	READ OR ACT? Action screen.	WHO AUTHORMANAGER	BULK UPLOAD Yes, several files at once.
---	---	---	---

What you can do

- 1

Upload. Drop in one file or many.
- 2

Rename. Renaming updates every template already pointing at the image, so nothing breaks.
- 3

Copy the public URL. For use outside Relay.
- 4

Download. Retrieve the original file.

The “unused” marker

An image no template currently references is marked `unused`. It is a housekeeping hint, not an instruction: an image staged for a campaign you have not written yet will look unused and should not be deleted.

✓ TIP

Size images before uploading. Email clients will not resize for you, and a very large file makes the message slow to open on a phone, which is where nearly everyone reads it.

Links

/links

VIEWER

AUTHOR

MANAGER

Relay's own short links, on the `lottomys.in` domain. They make a link short enough for SMS, let clicks be counted, and let a destination be changed after the message has gone.

WHAT IT'S FOR	READ OR ACT?	WHO	DOMAIN
Short, trackable, redirectable links.	Action screen.	AUTHOR MANAGER	<code>lottomys.in</code> , fixed.

Two kinds of link, with opposite rules

	RECIPIENT LINK	CAMPAIGN LINK
Code	Generated automatically	You choose the wording
Expires	After 30 days	Never
Destination	Fixed once created	Editable at any time
Used for	One customer's cart or order	Printed material, QR codes, a campaign

The difference is deliberate. A recipient link points at one person's cart, so it must never be repointed. A campaign link may be printed on a box that will exist for years, so it must never expire and must be repointable when the landing page moves.

⚠ IMPORTANT

Once a campaign link is printed on packaging or approved inside a WhatsApp template, the code cannot be taken back. Choose the wording carefully the first time. Changing where it points is fine and expected; changing the code itself is not possible after it is in the world.

Retiring a link

Retiring stops a link being offered for new sends. It does not switch off links already printed or already sent, which continue to work.

✓ **TIP**

The `?s=qr` marker distinguishes a scan of a printed code from a click in a message, so you can tell which channel actually produced the visit.

Analytics

/analytics

VIEWER

AUTHOR

MANAGER

APPROVER

ADMIN

SUPER-ADMIN

What happened after you sent. Delivery, engagement, cost and revenue, for a date range you choose.

WHAT IT'S FOR Measuring sends, honestly.	READ OR ACT? Read-only.	WHO Anyone with Relay access.	RANGE Presets, or a custom window. Defaults to today.
---	----------------------------	----------------------------------	--

The headline numbers

Delivery rate	Of what we sent, how much arrived.
Read rate	Of what arrived, how much was opened or read.
Our defects	Failures caused by us: a bad address, a variable that would not fill, a broken template. This is the number you can actually fix.
Throttled by Meta	Failures caused by WhatsApp rate limits rather than anything we did wrong. Kept separate on purpose so it is not read as a defect.
Skipped	Stopped by the send gate before going out, broken down by reason. These people received nothing.
Attributed revenue	Orders traced back to a send.
Marketing ROI	Attributed revenue against send cost. Transactional messages are excluded from this, because an order confirmation does not cause the order.

Where the rupee figures come from

Under the money tiles there is a **Cost basis · rate card** panel. It lists the rate charged for each kind of message, and the date each rate came into force.

Read it whenever a cost or ROI number looks wrong. The most likely explanation is not a bug — it is that the provider changed its prices and the rate card has not been updated yet, in which case every figure

above it is quietly out of date. **The *Effective from* column is what tells you.** A date that has not moved in a long time is the thing to question.

Spend is a floor, not a final number. A message is only costed once the provider tells us what it charged. A message that failed was never charged for, and one still in flight has not been priced yet — both are left out of Spend rather than counted as zero. So the real spend is never lower than the figure shown, and may be slightly higher.

SMS and RCS show no cost on purpose. Those providers report a credit balance rather than a per-message charge, and the two are not the same unit. Rather than show a plausible-but-wrong rupee figure, they are left out until the credits are reconciled.

Rates cannot be edited from this screen. Changing one re-prices *history* as well as new sends — every past message is costed at the rate that applied on the day it went out — so a repricing is done as a reviewed change to the database, not from a dashboard.

Sent, delivered, skipped, failed

These are four different things and mixing them up is the most common misreading on this screen.

STATUS	MEANS
sent	Handed to the provider.
delivered	Confirmed as arrived on the customer’s device.
skipped	The gate stopped it. Nothing was sent and nothing was charged.
failed	We tried and it did not work.

⚠ WATCH OUT

A large `skipped` count is not a system fault. It usually means the audience was right but the gate did its job: no consent, cap already hit, or outside quiet hours. Read the breakdown by reason before treating it as a bug, and remember those people were not contacted at all.

📌 NOTE

Email open tracking depends on the customer’s mail app loading images, which many block, so email read rates are systematically understated. WhatsApp read receipts are reliable. Compare email to email and WhatsApp to WhatsApp, never one against the other.

PART 6

Admin

Access, sending identities, and the safety limits that apply to everyone.

Users

/admin/users ADMIN SUPER-ADMIN

Who can open Relay, and what they can do once they are in.

WHAT IT'S FOR	READ OR ACT?	WHO	SCOPE
Granting, changing and switching off access.	Action screen.	ADMIN SUPER-ADMIN	Relay only. No other LOT system is affected.

Granting access

- 1 Grant access.** Find the person by their Legend of Toys account.
- 2 Pick a role.** Start with the smallest one that lets them do their job. Author is the right default for a marketer.
- 3 Confirm.** They can sign in immediately.

Active and inactive

Every grant carries an `active` switch. Setting someone inactive blocks them from Relay at once while keeping the record of what they had. Use it when someone changes team or leaves: it is instant, reversible, and leaves an audit trail that deleting does not.

⚠ WATCH OUT

Relay access is separate from every other system. Removing someone here does not remove them from Garage, Pitstop or anywhere else, and offboarding somebody properly means doing it in each system they had.

📌 NOTE

Some accounts in the list are service principals rather than people, used by internal tooling to read data. They look like odd non-human entries and should be left alone: removing one silently breaks whatever depends on it.

Roles

/admin/roles

SUPER-ADMIN

The permission builder. Defines what each role can do, for everyone in Relay.

WHAT IT'S FOR	READ OR ACT?	WHO	APPLIES
Editing the roles themselves.	Action screen, with system-wide effect.	SUPER-ADMIN only.	Immediately, to everyone holding that role.

The permission keys

KEY	ALLOWS
relay_view	Opening Relay and reading screens. Without this, nothing.
segment_manage	Creating and editing segments.
template_manage	Creating and editing templates.
campaign_build	Preparing campaigns and journeys, but not releasing them.
send_activate	Pressing send, scheduling, and activating a journey.
approve	Approving a campaign over the threshold.
data_consent_admin	Editing consent and delivery blocks on a contact.
connector_channel_manage	Managing sender identities and connectors.
relay_admin	Granting access to others.
relay_super_admin	This screen, and Approval & Caps.

 IMPORTANT

Keep `campaign_build` and `send_activate` apart in any role meant for someone learning the system. Granting both is what turns “I was just trying it out” into a live send to twenty thousand people.

① **NOTE**

Roles marked `system` are read-only and cannot be edited or deleted. They are the safe baseline. Build a new role rather than trying to bend one of those.

Approval & Caps

The safety limits that apply to every send in Relay, regardless of who made it. The highest-consequence screen in the app.

WHAT IT'S FOR	READ OR ACT?	WHO	APPLIES
Test mode, approval threshold, frequency cap, quiet hours, send budget.	Action screen, system-wide.	SUPER-ADMIN only.	Within a minute, to everything.

Test mode: the global send lock

While test mode is on, no message reaches anyone outside the internal test list. Campaigns still run and still report, but nothing lands on a customer. It is the master safety switch and the thing to check first whenever a send appears to have gone nowhere.

🚫 IMPORTANT

Switching test mode off arms the entire system at once, including every active journey. Do it deliberately, tell the team, and check the Overview screen afterwards to confirm the state you intended.

The two allowlists

Two separate lists decide who a test send may reach, and they are edited in the same place but owned by different people.

Allowlist	The policy list. Accepts a whole domain with a leading @, or an exact email. While test mode is on, only these addresses receive anything.
Test-send allowlist	Exact emails and phone numbers only — no @domain patterns. The Send-test buttons on Templates and Campaigns add to this list automatically, so it grows on its own as people test.

A test send may only ever reach someone on one of these two lists, **even when test mode is off**. That is deliberate: a test send carries arbitrary, unapproved content.

⚠ WATCH OUT

Removing a number from the Test-send allowlist stops that person's test sends working immediately. If a colleague reports "my test send is refused", check here first — the same number written two ways (`9876543210` and `+919876543210`) counts as two different entries, and only the exact form used in the Send-test box will match.

Send governance

Approval threshold	Marketing campaigns above this audience size need an Approver. Below it, they approve themselves. Judged on the number that will actually receive, after exclusions.
Frequency cap	The maximum marketing messages one person can get in a rolling window. It is a rolling window, not a calendar day.
Daily send budget	A ceiling on marketing sends per day, for warming up a sending number.

⚠ WATCH OUT

If the budget is lower than the audience, the campaign will still report as sent while quietly skipping everyone past the limit, with the reason `budget_exhausted` . Those people are not queued and will not be retried. Set the budget above the audience before a large broadcast.

Quiet hours, per channel

Set separately for each channel because the constraints genuinely differ: promotional SMS in India is legally restricted, WhatsApp is not, and email is exempt. A channel can be switched off from quiet hours entirely. Marketing and survey messages respect quiet hours; order updates do not.

Link tracking defaults

Account-wide UTM defaults. A template or campaign can override any field, and a blank field inherits rather than sending an empty value.

Sender Identities

[/admin/senders](#)

ADMIN

SUPER-ADMIN

The addresses and numbers Relay sends from. Which one a message uses is decided here, not on the campaign.

WHAT IT'S FOR Managing the from address and sending numbers.	READ OR ACT? Action screen.	WHO ADMIN SUPER-ADMIN	CHANNELS Email, WhatsApp, SMS.
---	--------------------------------	--------------------------	-----------------------------------

What a sender identity is

One row per address or number Relay may send from, with its channel and whether it is active. Only an active identity is used. Marking one inactive takes it out of rotation without deleting the history attached to it.

Why more than one number

Separating marketing traffic from order updates protects the important messages. If a marketing number gets rate limited or its quality rating drops, order confirmations carry on unaffected because they leave from somewhere else.

⚠ WATCH OUT

Changing the email from address affects deliverability. A new domain has no sending reputation and will land in spam more often until it builds one. Do not swap it casually before a large send.

📌 NOTE

On WhatsApp, a template belongs to a specific business account, which constrains which numbers can send it. If a template will not go out from the number you expected, this is usually why.

Connectors

/admin/connectors

ADMIN

SUPER-ADMIN

The feeds that bring outside data into Relay. When something stops arriving, this is where you look.

WHAT IT'S FOR Watching the integrations Relay depends on.	READ OR ACT? Mostly read.	WHO ADMIN SUPER-ADMIN	CREDENTIALS Not here. They are held securely outside the app.
--	------------------------------	--	--

The connectors

Shopify	Customer data: who exists, what they bought, what they are worth. This is what fills the attributes segments filter on.
Cashfree	Payment links, used where a message asks a customer to pay.

When a feed stops

A stalled feed is quiet rather than loud. Segments keep returning their last answer and journeys simply stop firing, with no error on any screen. The reliable early warning is the Activity screen: if a familiar event has gone to zero today, come here.

① NOTE

Not everything arrives from a connector. Some data is pushed to Relay by other Legend of Toys systems as it happens. A connector looking healthy therefore does not prove every feed is healthy, and Activity remains the better overall check.

⚠ WATCH OUT

Marketing consent is one field Shopify does not currently pass through. A customer's consent in Relay comes from their interactions with us, not from their Shopify record, so do not treat a Shopify opt-in as the truth here.

Suppressions

/suppressions

ADMIN

SUPER-ADMIN

Everyone Relay will not send to, and why. A suppression is the hardest block in the system — it stops every message on that channel, order updates included.

WHAT IT'S FOR	READ OR ACT?	WHO	CHANNELS
Seeing who is blocked, and lifting a block when it is wrong.	Mostly a read screen. One action: lift.	CONSENT ADMIN	Email, WhatsApp, SMS.

What a suppression actually does

It is the first check in the send gate, and it is absolute. A suppressed address or number receives nothing on that channel — not marketing, not order confirmations, not shipping updates. This is what makes it different from an opt-out, which only stops marketing.

That is also why this screen exists. When someone asks “why did this customer never get their order confirmation?”, a suppression is the most common answer, and until now the only way to check was to ask an engineer.

The reasons, and what they mean

Reason	Where it comes from	Lift it?
hard_bounce	The mail provider told us the address does not exist.	Only if you know the address is genuinely valid.
complaint	The recipient marked one of our emails as spam.	Rarely. See the warning below.
wa_undeliverable_...	WhatsApp rejected delivery three times in a row.	Usually pointless — see below.
gdpr_redact	A legal erasure request.	Never. The button is not offered.
manual	Somebody added it by hand.	Yes, if it was a mistake.

A long list here is normal

The first reaction to seeing thousands of rows is that something is broken. For the largest group, the opposite is true.

A WhatsApp number is suppressed only on its **third** consecutive hard failure. That threshold is deliberate: roughly one in thirteen numbers that fail once will go on to deliver perfectly well later, so

suppressing after a single failure would quietly cut off reachable customers to save a fraction of a message. The retries are how the genuinely dead numbers get found.

So a count that grows is the mechanism doing its job. The only figure worth investigating is one that *stops falling* over a period of months.

⚠ WATCH OUT

Lifting a `complaint` block is the riskiest thing on this screen. That person told their mail provider our message was spam. Sending to them again risks the sender reputation that every other customer's email depends on — including order confirmations.

Lifting a block

Find the row, press **Lift block**, and confirm. The block is removed immediately and the person becomes reachable again on that channel.

Every lift is recorded against your name in the **Lifted** tab, with what the block was and when it was removed. That tab is the *only* record that the block ever existed — lifting deletes the original row — so it is worth a look before you assume a customer was never blocked at all.

① NOTE

There is no “clear all”, and that is deliberate. Bulk-lifting would re-admit thousands of known-dead addresses to every future send in one click, undoing the mechanism described above. Blocks come off one at a time, with a name against each.

Finding one person

Search matches on the email address or phone number. You can also narrow by channel, or by an exact reason. If you are looking at a single customer, the **Contacts** screen shows the same blocks on their record alongside everything else about them — that is usually the faster route.

Exporting

Export CSV downloads every row matching the filters you have applied, not just the page on screen. If the export is ever truncated you are told so and asked to narrow the filters — a short file is never handed over as if it were complete.